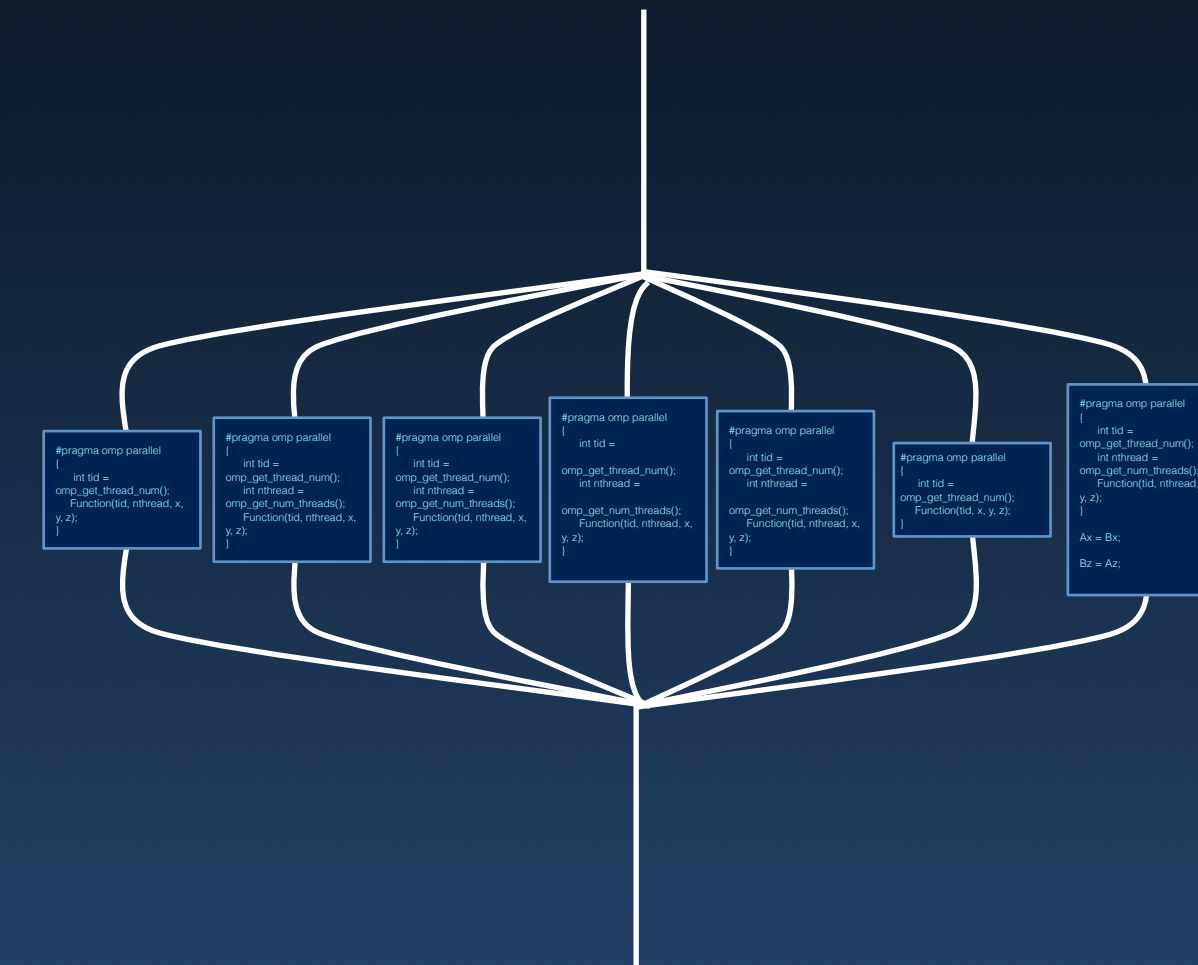


COL380

Introduction to
Parallel & Distributed Programming

Hello OpenMP

- Fork-Join programming model
 - ➔ Teams of threads
- Shared Memory programming model
- Well defined memory model
- High level synchronization

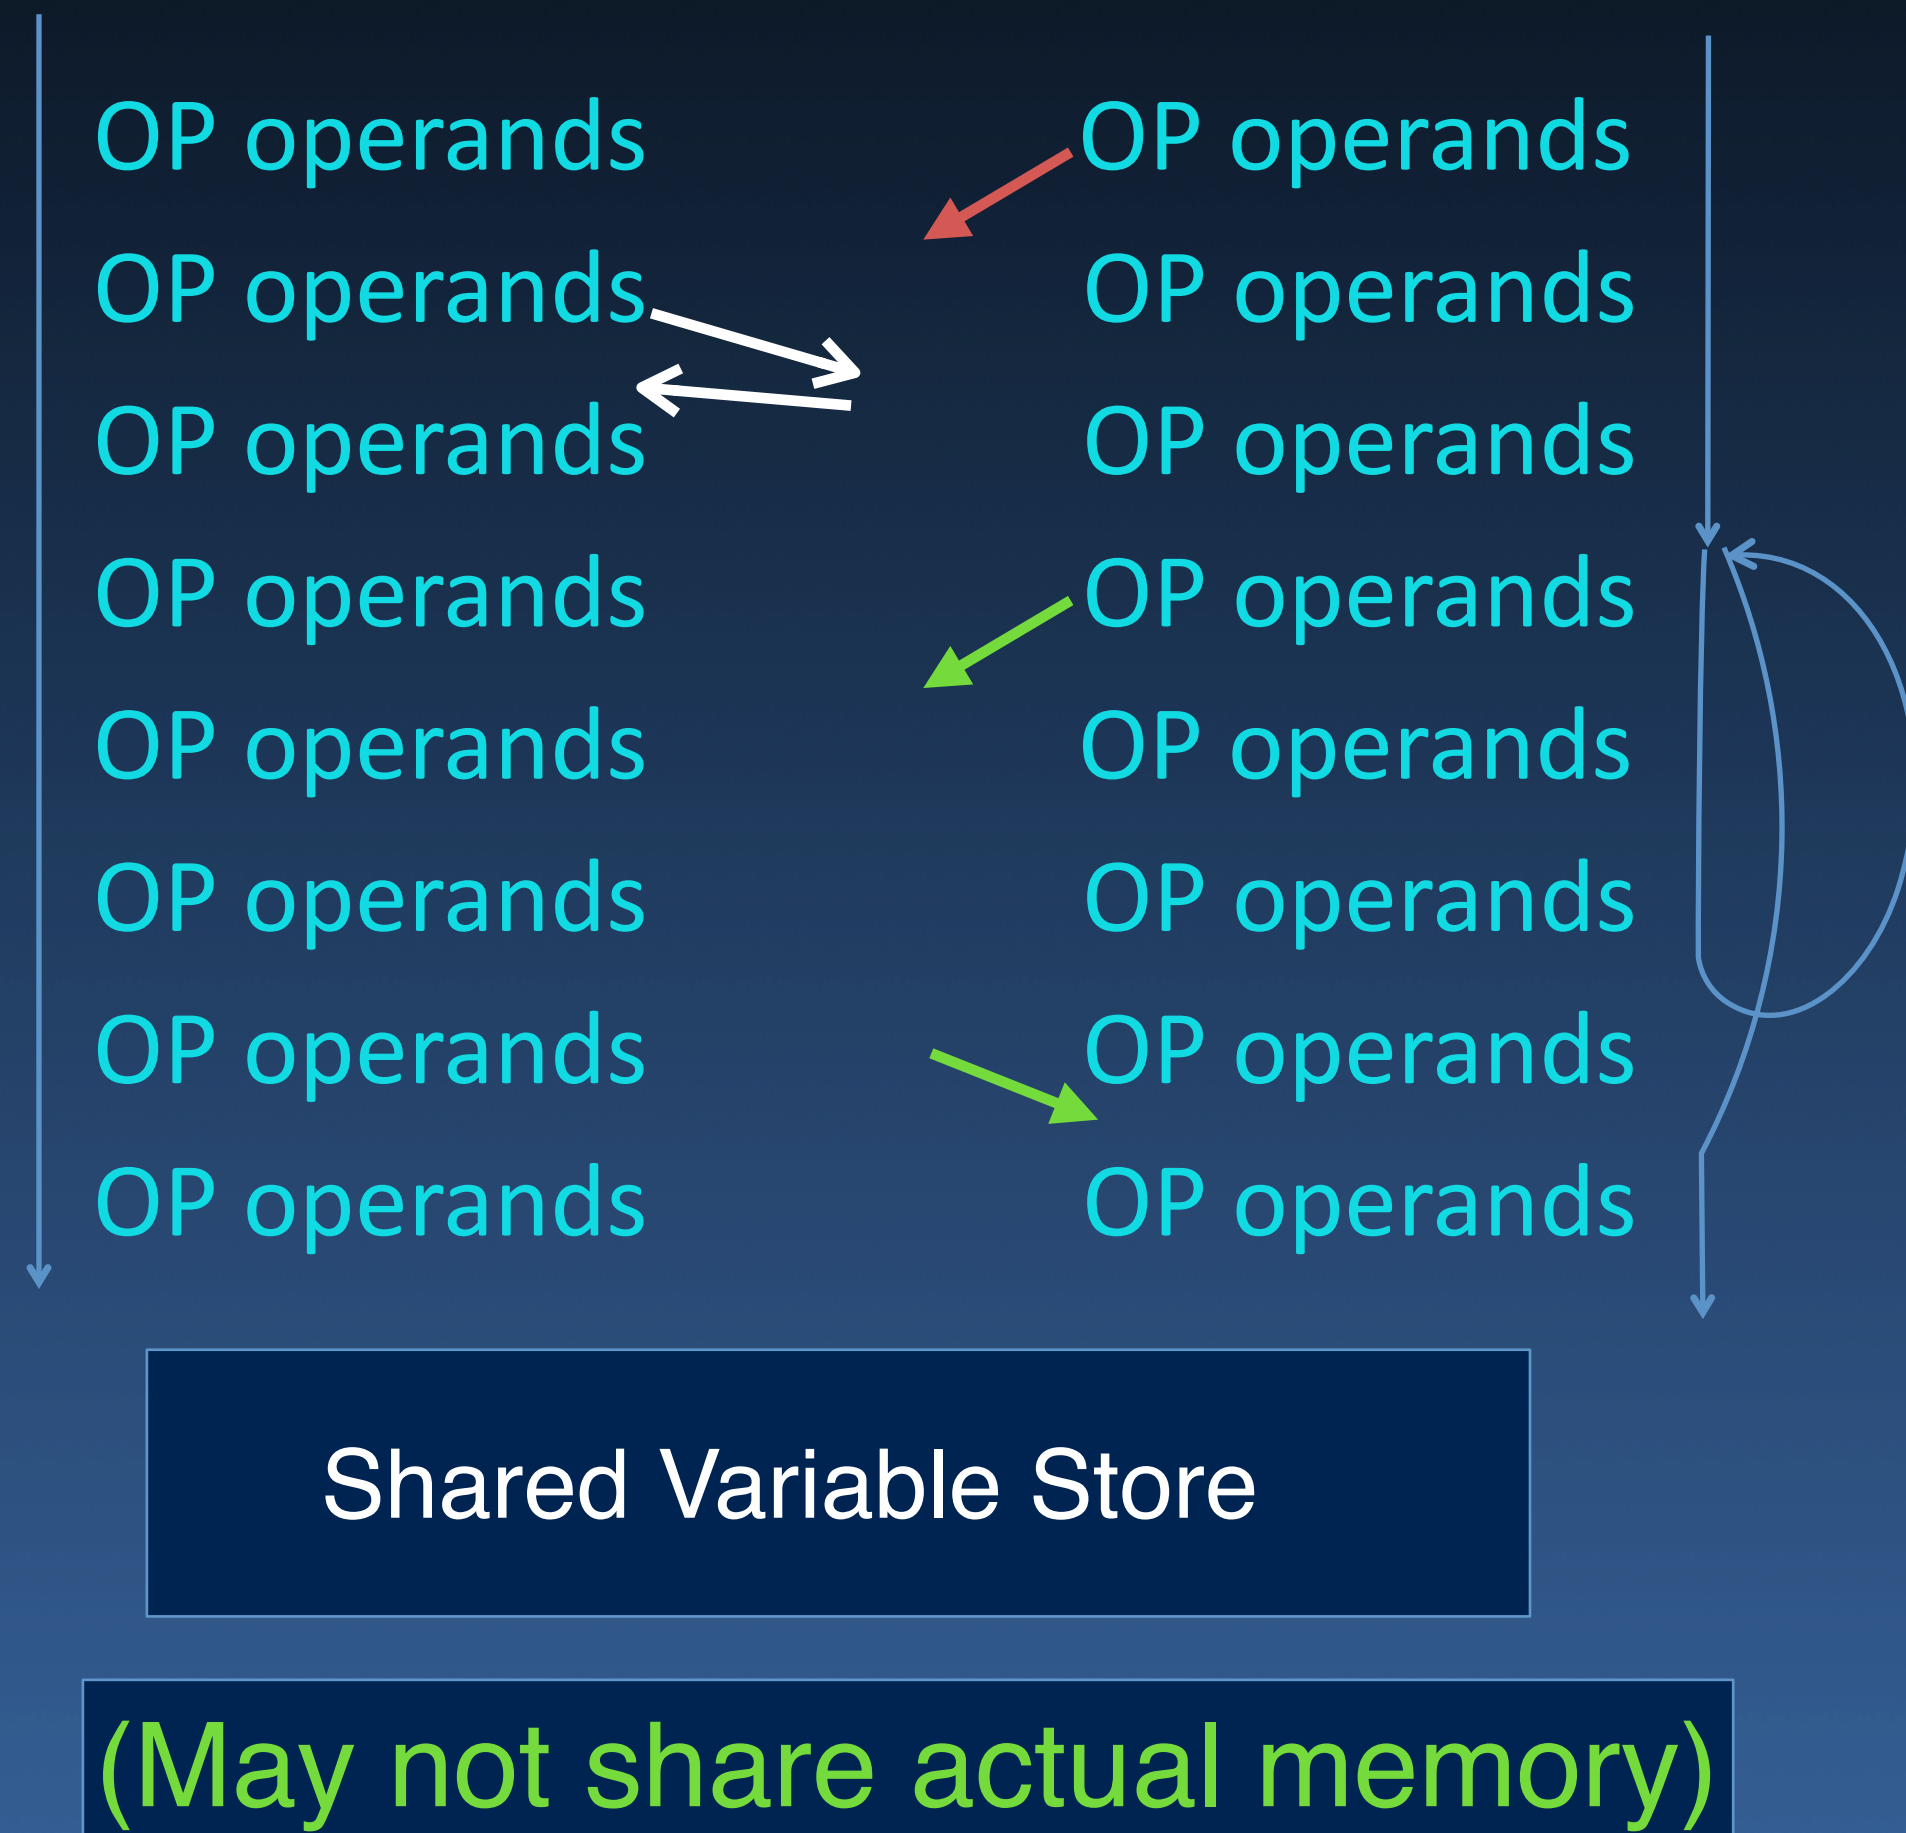


```
#pragma omp parallel
{
    int tid = omp_get_thread_num();
    int nthread = omp_get_num_threads();
    Function(tid, nthread, x, y, z);
}
```

- “Light-weight” process
 - ➔ Execution context: Stack, Registers
 - ➔ Context switch of Registers, SP, PC
 - ▶ Not caches, virtual page tables
- **Kernel threads vs User threads** (Light-weight thread)
 - ➔ Kernel threads are scheduled by OS
 - ➔ User threads when no context switching is required

Parallel Program

Multiple Threads of Execution



- How are they instantiated?
- Execution and memory model
- How do they interact?

Separate from HW
(Who does the execution?)

Programming Models

- Shared Memory model
- Distributed Memory model
- Task based model
- Work-queue model
- Stream processing model
- Map-reduce model
- Client-server model

Example Models

```
int threadFn(int arg)
{
    // Do something with 'arg'
}
...
{
    int x;

    x = spawn threadFn(1);
    // Continue doing other things
    sync;
}
```

```
void *threadFn(void *arg)
{
    // Do something with 'arg'
}
...
{
    int arg;
    pthread_t thread_id;
    pthread_create(&thread_id, NULL, threadFn, arg);
    // Continue doing other things
    pthread_join(thread_id, NULL);
}
```

```
func threadFn(id int, arg string, wg *sync.WaitGroup,
              ch chan int)
{
    defer wg.Done()
    // Do thing 'id' with 'arg'
    x <- ch
}
...
{
    var wg sync.WaitGroup
    chA := make(chan int)

    wg.Add(1)
    go threadFn(1, "a string", &wg, chA)
    // Continue doing other things
    wg.Wait()
    chA <- result
}
```

- High level language

```
for i=0 to N  
  a[i] = f(b[i], c[i], d[i])
```

- ➔ Derive parallelism
- ➔ Generate threads and map to processors
- ➔ Addresses for a, b, c, d accessible to all
 - ▶ also the code for f
- ➔ Map i to thread ID
- ➔ Impact on cache coherence?

- **Generate new threads of control**
 - ➔ function per thread, work sharing construct
- **Synchronize**
- **Specify variable scopes**
 - ➔ Maybe, for an arbitrary group of threads
- **Ways to map threads to processor?**
 - ➔ May have more threads than processors
 - ➔ Scheduling hints

- Encountering thread creates a team:
 - ➔ Itself (master) + zero or more additional threads.
- Applies to structured block immediately following
 - ➔ Each thread executes separately the code in {
 ▶ But, also see: Work-sharing constructs
- There's an implicit barrier at the end of block
- Only master continues beyond the barrier
- May be nested
 - ➔ disabled by default

- Notion of temporary view of memory
 - ➔ Allows local caching
 - ➔ Need to relax consistency model
- Supports threadprivate memory
 - ➔ global scope
- Variables declared before parallel construct:
 - ➔ Shared by default
 - ➔ May be designated as private
 - ▶ $n-1$ copies of the original variable is created

Initialized if requested

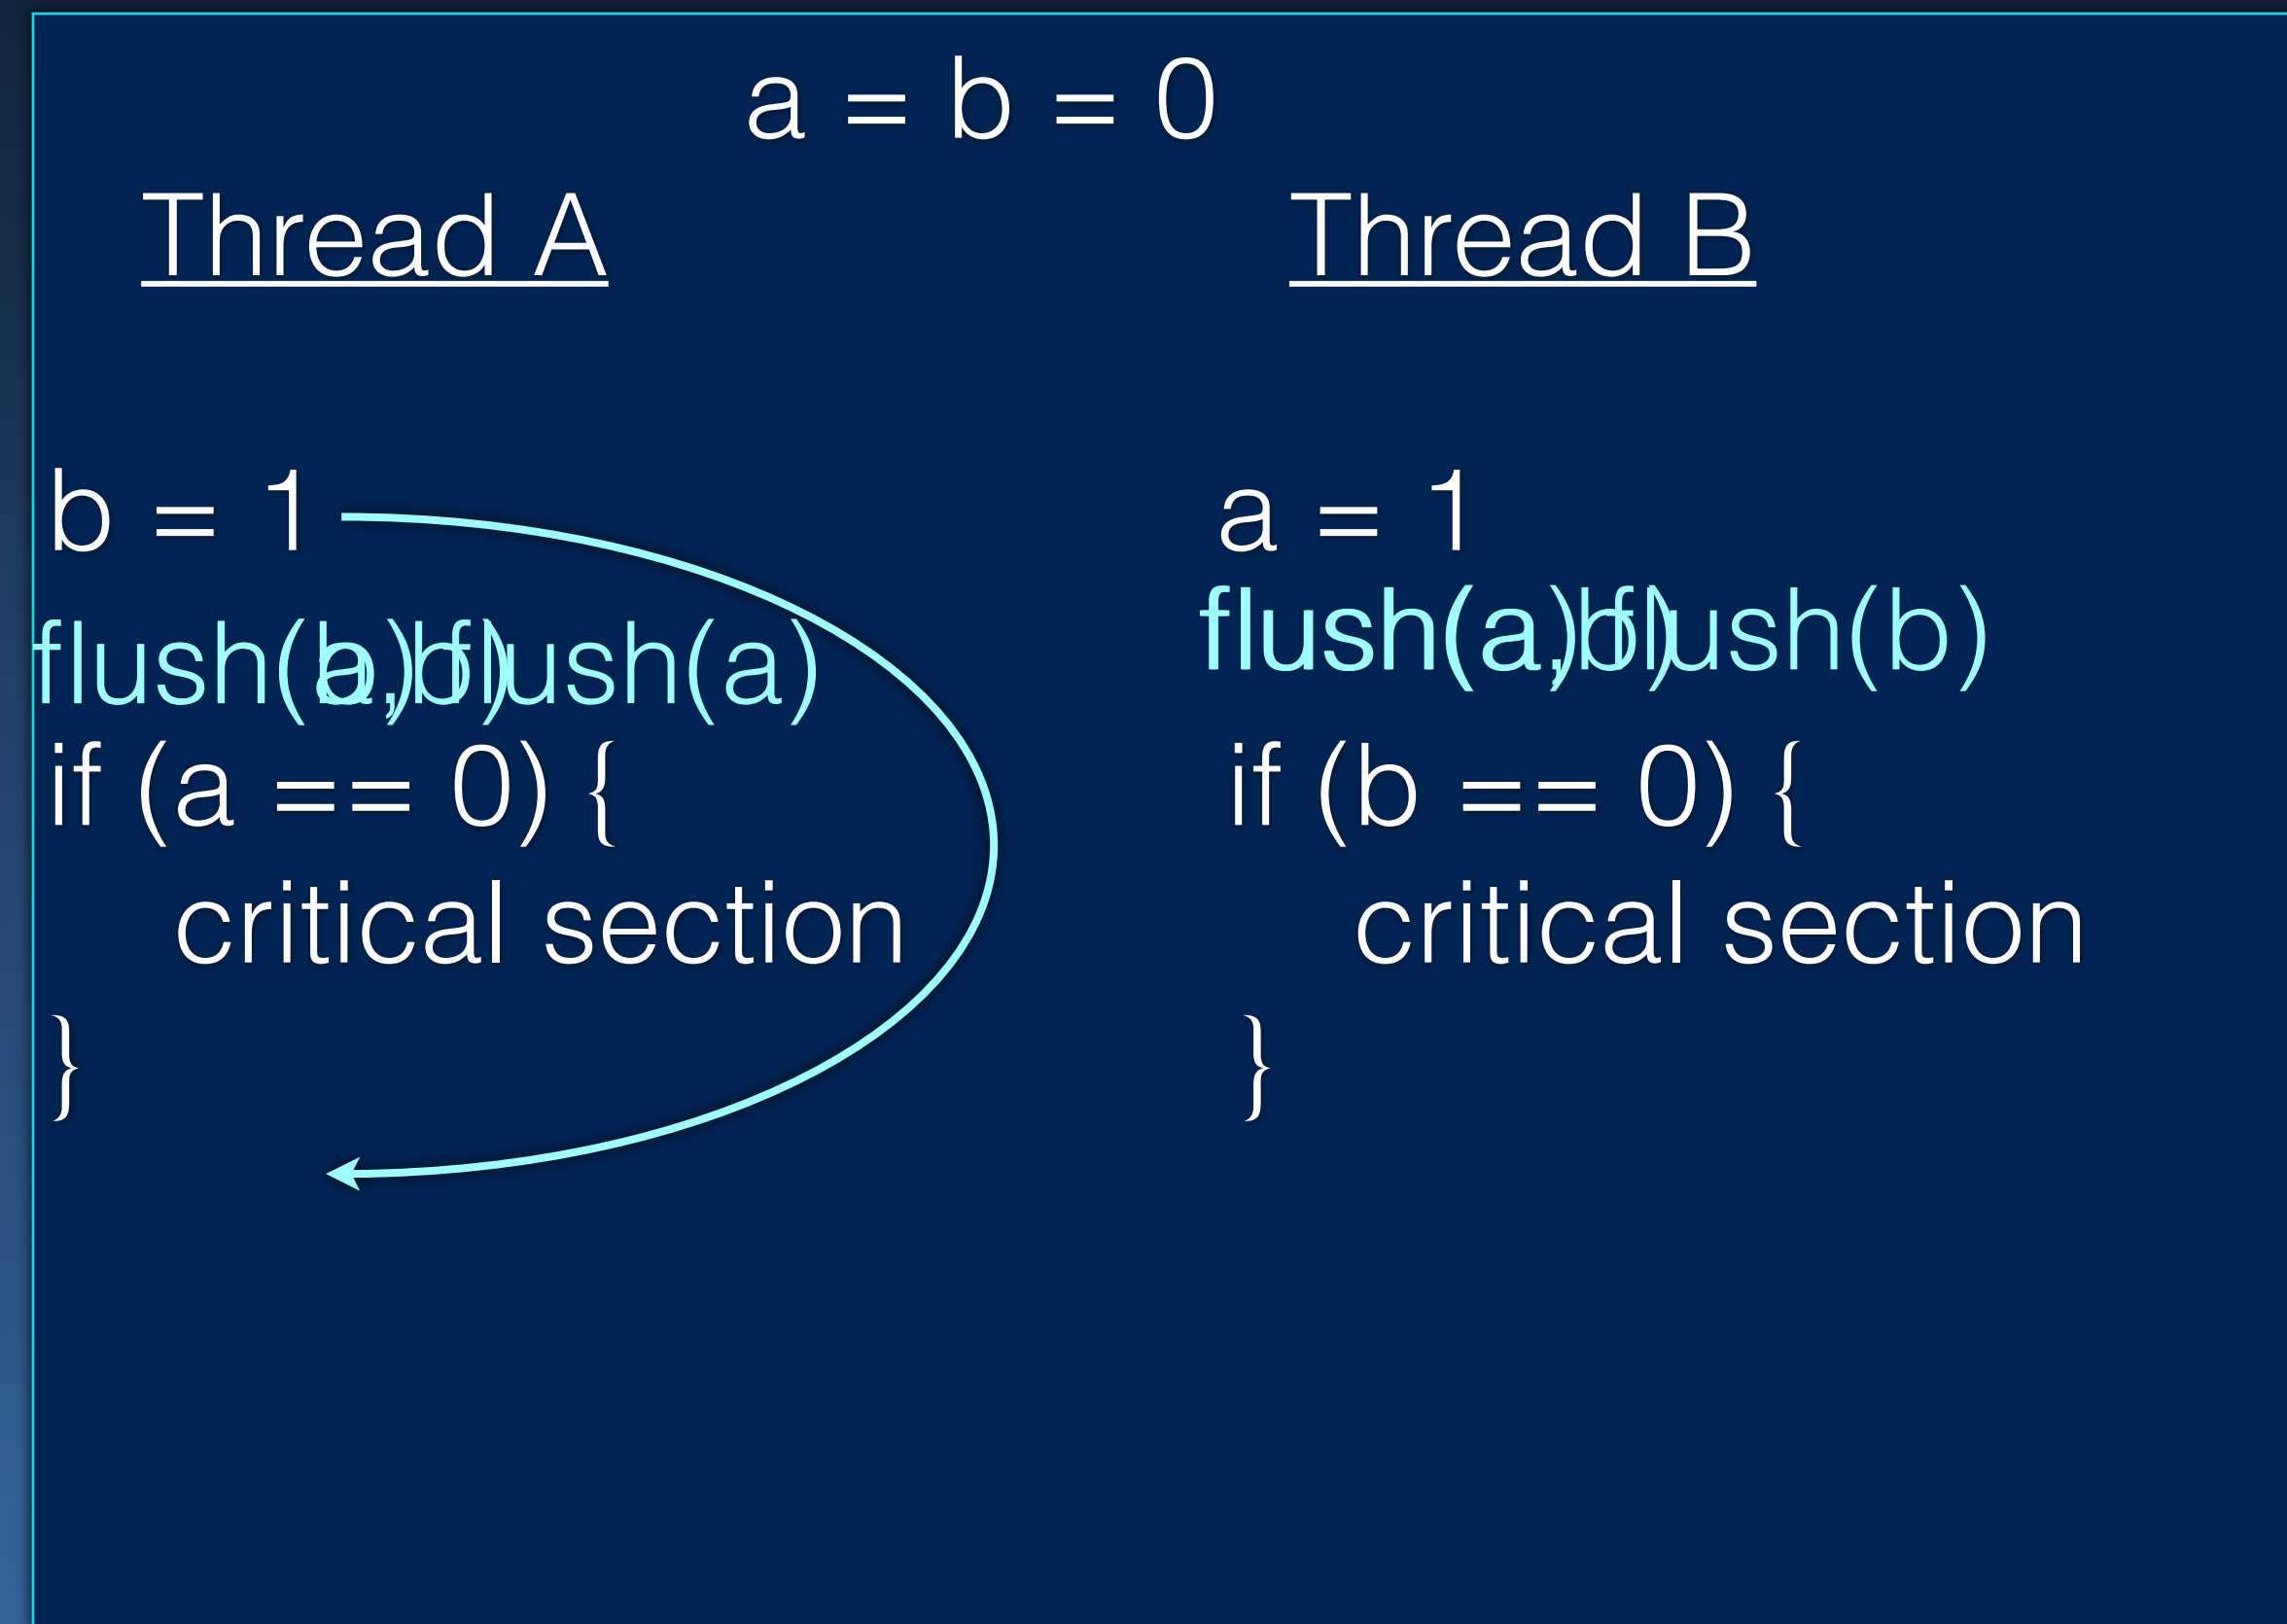
Sharing Variables

- **Shared:**
 - ➔ Heap allocated storage
 - ➔ Static data members, const variable (no mutable member)
- **Private:**
 - ➔ auto Variables declared in a scope inside the construct
 - ➔ Loop variable in for construct
- **Others are shared by default** Can declare private or change default
- **Arguments passed by reference inherit from original**

- **Unsynchronized access:**
 - ➔ If two threads write to the same shared variable the result is undefined
 - ➔ If a thread reads and another writes, the read value is undefined
- Memory atom size is implementation dependent
- Flush x,y,z .. enforces consistency

<u>Thread A</u>	<u>Thread B</u>
X = 5	Flush(X)
...	...
Flush(X)	v = X

Compiler Re-ordering



`#pragma omp flush(a,b)`

1-Producer 1-Consumer

int data, flag = 0;

Thread 0

```
// Produce data
data = 42;
// Flush
F #pragma omp flush(flag, data)
// Set flag to signal Thread 1
A flag = 1;
// Flush
#pragma omp flush(flag)
```

produce

Thread 1

```
// Busy-wait until flag is signalled
#pragma omp flush(flag)
while (flag != 1) {
    #pragma omp flush(flag) data
}
#pragma omp flush(flag, data)
// Consume data
printf(data="%d\n", data);
consume
```

“1” ⇒ A has started and hence F has happened

Acquire & Release Flush

$a = b = 0$

Thread A

$b = 1$

#pragma omp flush release

if ($a == 0$) {

critical section

}

Thread B

$a = 1$

#pragma omp flush acquire

if ($b == 0$) {

critical section

}