

COL380

Introduction to
Parallel & Distributed Programming

Parallel Construct

```
#pragma omp parallel \  
    if(boolean) \  
    private(var1, var2, var3) \  
    firstprivate(var1, var2, var3) \  
    default(private | shared | none) \  
    shared(var1, var2) \  
    copyin(var1, var2) \  
    reduction(operator:list) \  
    num_threads(n)  
{  
}
```

- Implicit barrier at the end, Implicit flush
- Cannot branch in or out
- No side effect from clause: must not depend on any ordering of the evaluations
- Upto one if and *num_threads* clauses
- num_threads must be a +ve integer

Variable Scope

```
int size;  
int numProcs = omp_get_num_procs();  
#pragma omp parallel num_threads(numProcs)  
{  
    size = getProblemSize()/numProcs;  
    int tid = omp_get_thread_num();  
    doTask(tid*size, size);  
}  
  
void doTask(int start, int count)  
{ // Each thread's instance has its own activation record  
    for(int i = 0, t=start; i< count; i++; t++)  
        doit(t);  
}
```

Private clause

```
int tid, size;
int numprocs = omp_get_num_procs();
#pragma omp parallel num_threads(numProcs) private(tid)
{
    size = getProblemSize()/numProcs;
    tid = omp_get_thread_num();
    doTask(tid*size, size);
}
```

```
void doTask(int start, int count)
{ // Each thread's instance has its own activation record
    for(int i = 0, t=start; i< count; i++; t++)
        doit(t);
}
```


- `#pragma omp flush (var1, var2)`
 - ➔ Stand-alone, like barrier
 - ➔ Only directly affects the encountering thread
 - ➔ List-of-vars ensures that any compiler re-ordering moves all flushes together
 - ➔ implicit:
 - ▶ barrier, atomic, critical, locks

```
#pragma omp parallel for  
  for (i= 0; i < N; i++) {  
    blah ...  
  }
```

- Num of iterations must be known when the construct is encountered
 - ➔ Must be the same for each encountering thread
- Compiler puts a barrier at the end of **parallel for**
 - ➔ But see **nowait**

Parallel For Construct

- `#pragma omp for \`

- ➔ `private(var1, var2, var3) \`

- ➔ `firstprivate(var1, var2, var3) \`

Once per thread, not once per iteration

Original not corrupted

- ➔ `lastprivate(var1, var2) \`

Last iteration's value

- ➔ `reduction(operator: list) \`

- ➔ `ordered \`

- ➔ `schedule(kind[,chunk_size]) \`

- ➔ `nowait`

- ➔ Canonical For Loop

No break

- same loop control expression for all threads in the team.
- At most one `schedule`, `nowait`, `ordered` clause
- `chunk_size` must be a loop/construct invariant, +ve integer
- `ordered` clause required if any `ordered` region inside

```
#pragma omp ordered
```

```
{  
}
```

- ➔ Binds to inner-most enclosing loop
- ➔ The structured block executed in loop sequential order
- ➔ The loop must declare the ordered clause
- ➔ Each thread must encounter only one ordered region

- Encountering thread creates a task
 - ➔ Code, data, environment ..
- ‘Some’ thread of the team executes the task
 - ➔ Scheduling points
 - ▶ Start, End, taskwait Barrier

Task Construct

**#pragma omp task **

if(boolean) \

untied \

default(shared | none) \

private(list) \

firstprivate(list) \

shared(list) \

depend(modifier:list)

{
}

- Cannot branch in or out
- No side effect from clause: must not depend on any ordering of the evaluations
- Upto one if clause

#pragma omp taskwait
#pragma omp taskgroup
#pragma omp taskyield

Traversal By Tasks

```
struct node { node *left, *right; Data *data;};  
extern void process(node* );  
void traverse(node* p)  
{  
    if (p->left)  
        #pragma omp task // p is firstprivate by default  
        traverse(p->left);  
    if (p->right)  
        #pragma omp task // p is firstprivate by default  
        traverse(p->right);  
    process(p); ← #pragma omp taskwait  
}
```

Recursive Sum

```
float seq_sum(const float *a, int n)
{
    return (n == 0)? 0.
        : (n == 1)? *a
        : seq_sum(a, n/2) +
          seq_sum(a + n/2, n - n/2);
}
```

```
#pragma omp parallel
#pragma omp single nowait
sum = par_sum(a, n);
```

```
float par_sum(const float *a, int n)
{
    if (n <= SMALL)
        return seq_sum(a, n);
    float x, y;

    #pragma omp task shared(x)
    x = parallel_sum(a, n/2);
    #pragma omp task shared(y)
    y = parallel_sum(a + n/2, n - n/2);
    #pragma omp taskwait
    x += y;

    return x;
}
```


- **reduction (<op>: <variable list>)**
 - ➔ + Sum
 - ➔ * Product
 - ➔ & Bitwise and
 - ➔ | Bitwise or
 - ➔ ^ Bitwise exclusive or
 - ➔ && Logical and
 - ➔ || Logical or
- **Add to parallel for**
 - ➔ OpenMP creates a loop to combine copies of the variable
 - ➔ The resulting loop may not be parallel

Reduction

```
int sumint(int data[]) {  
    int sum;  
    #pragma omp parallel reduction(+:sum)  
        sum = partial_sum(data, omp_get_thread_num());  
    return sum;  
}
```

Reduction

```
int sum2d(int data[N][N]) {  
    int sum;  
    #pragma omp parallel for reduction(+:sum)  
    for (int i=0; i<N; i++) {  
        for (int j=0; j<N; j++) {  
            sum += data[i][j];  
        }  
    }  
    return sum;  
}
```

```
#pragma omp critical (accessBankBalance)
{
}
```

- A single thread at a time
 - ▶ through all regions of the same name
- Applies to all threads
- The name is optional
 - ▶ Anonymous = global critical region

#pragma omp barrier

- Stand-alone
- Binds to inner-most parallel region
- All threads in the team must execute
 - ▶ they will all wait for each other at this instruction
 - ▶ Dangerous:
if (! ready)
#pragma omp barrier
- Same sequence of work-sharing and barrier for the entire team

`#pragma omp atomic`

`x++;`

read/write/update/capture

- Light-weight critical section

- Only for some basic operations

→ `x binop= expr` (no

→ `x++`

→ `--x`

→ `v = x++`

Thread 0

// Produce data
`data = 42;`

// Set flag to signal Thread 1
`#pragma omp atomic write`
`flag = 1;`

Thread 1

// Busy-wait until flag is signalled
`#pragma omp atomic read`

`myflag = flag`
`while (myflag != 1) {`
`#pragma omp atomic read`
`myflag = flag`
`}`

// Consume data
`printf("data=%d\n", data);`

Atomic?

```
int sum2d(int data[N][N]) {  
    int sum = 0;  
    #pragma omp parallel for  
    for (int i=0; i<N; i++) {  
        for (int j=0; j<N; j++) {  
            #pragma omp atomic  
            sum += data[i][j];  
        }  
    }  
    return sum;  
}
```

- Error due to non-deterministic ordering of mutually visible event

- Data races:

- RW

- ▶ Reading of an object potentially overlapped

- WW

- ▶ Potentially overlapping writes of an object

```
Transfer(from, to, amt)
{
    #pragma omp atomic read
    fbal = from.balance;
    if (fbal < amt) return ERROR;
    #pragma omp atomic update
    to.balance += amount;
    #pragma omp atomic update
    from.balance -= amount;
    return OK;
}
```


- Lock is an abstract datatype, supports three operations:
- **new** creates a new lock
 - ➔ Initially “open” by default
- **acquire** “closes” a lock if open
 - ➔ blocks if closed, until someone else opens it
- **release** opens a lock

Re-entrant Lock

- **new** creates a new lock with no current holder and a count of 0
- **acquire** blocks *if held by someone different* from the caller
 - ➔ If caller is the holder, increment count
 - ➔ If not already held by a different holder, caller gets a hold with count = 1
- **release** sets the current holder to “none” if the count is 0.
 - ➔ Otherwise, decrement count

Helper Functions

- `void omp_init_lock (omp_lock_t *);`
 - ➔ `void omp_destroy_lock (omp_lock_t *);`
- `void omp_set_lock (omp_lock_t *);`
- `void omp_unset_lock (omp_lock_t *);`
- `int omp_test_lock (omp_lock_t *);`
- **nested lock versions:**
 - ➔ e.g., `omp_set_nest_lock(omp_test_lock_t *);`

```
DoX()
{
    lock(X)
    Operate Exclusively()
    unlock(X)
}
```

```
DoXPlus
{
    lock(X)
    if(not setup X)
        Do(X)
    unlock(X)
}
```

- **Minimize synchronization**
 - ➔ Avoid BARRIER, CRITICAL, ORDERED, and locks
 - ➔ Use NOWAIT
 - ➔ Use named CRITICAL sections for fine-grained locking
 - ➔ Use MASTER (instead of SINGLE)
- **Parallelize at the highest level possible**
 - ➔ such as outer FOR loops
 - ➔ keep parallel regions/tasks large

- FLUSH is expensive
- LASTPRIVATE has synchronization overhead
- Thread safe malloc/free are expensive
- Reduce False sharing
 - ➔ Careful design of data structures
 - ➔ Use PRIVATE

- Try to
 - ▶ Avoid nested locks
 - ▶ Release locks religiously
 - ▶ Avoid “while true” (especially, during testing)
- Be careful with
 - ➔ Non thread-safe libraries
 - ➔ Concurrent access to shared data
 - ➔ IO inside parallel regions
 - ➔ Differing views of shared memory (FLUSH)
 - ➔ NOWAIT