


COL380

Introduction to
Parallel & Distributed Programming

Wait free

Registers

- Regular



- Single writer

Quiescently consistent

- Reads overlapping with a write return the newly written value or the one prior

- Atomic

A later read cannot see an earlier write

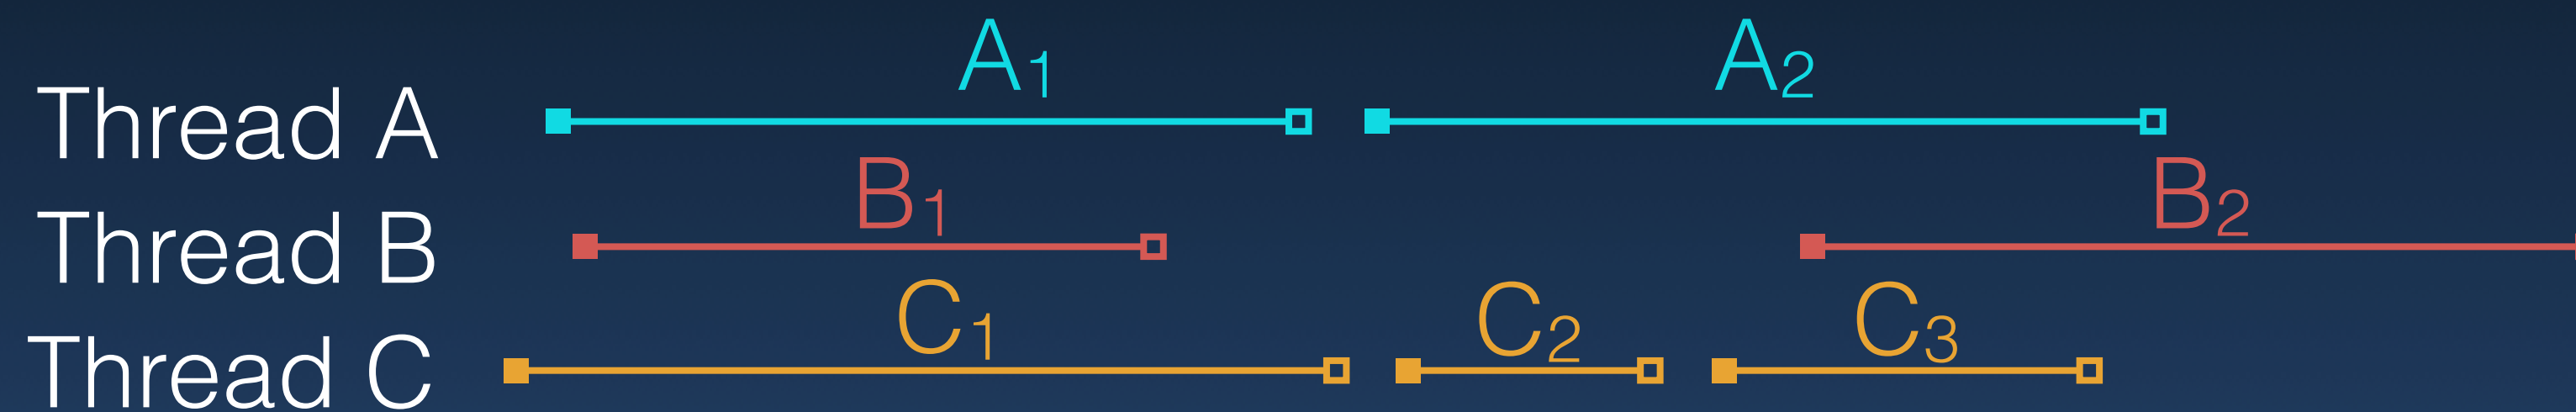
- Multiple writers

Linearizable

- Read the value of latest write (*takes effect* at some point in its duration)

Linearizable

- Operations appear to take effect at some instant
→ between their start and end



Equivalent sequential history S exists:

A₁ C₁ B₁ C₂ B₂ C₃ A₂

1. Each thread's history is retained in S
2. Non 'overlapping' operations retain order

Composable

Strict Consistency

- Read(x) returns the value stored by the 'most recent' write(x)
 - ➔ must appear as if operations are completed instantaneously
 - ➔ requires knowledge of absolute global time
 - ➔ intuitive to reason in this model
- Linearizability is weaker
 - ➔ accesses have a linearization point between their start and completion
 - ▶ and appear to occur instantly at the linearization point
 - ➔ Need a notion of a global clock
 - ➔ Useful for reasoning

- “A multiprocessor is **sequentially consistent** if the result of any execution is the same as if the operations of all the processors were executed in **some sequential order**, and the operations of each individual processor appear in this sequence **in the order specified by its program**.” [Lamport, 1979]

< Weaker than Linearizability


Sequentially Consistent

Thread A: $\frac{X = 5}{\text{EnQ}(5)}$


Thread B: $\frac{X = 3}{\text{EnQ}(3)}$ $\frac{\text{Read } X (3)}{\text{DeQ is 3}}$ $\frac{\text{Read } X (5)}{\text{DeQ is 5}}$

- No global notion of time

→ Only consistent Order

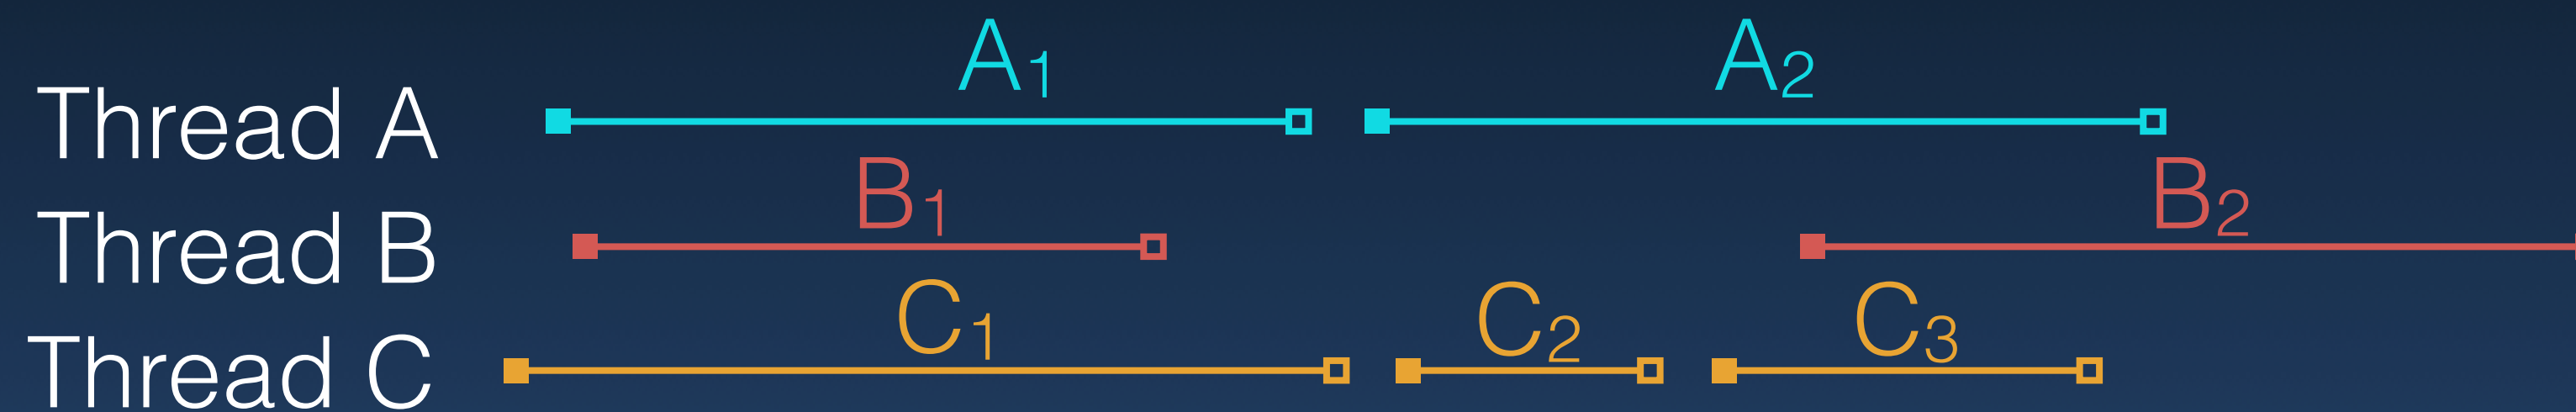


Sequential History



Linearizable

- Operations appear to take effect at some instant
→ between their start and end



Equivalent sequential history S exists:

A₁, C₁, B₁, C₂, B₂, C₃, A₂

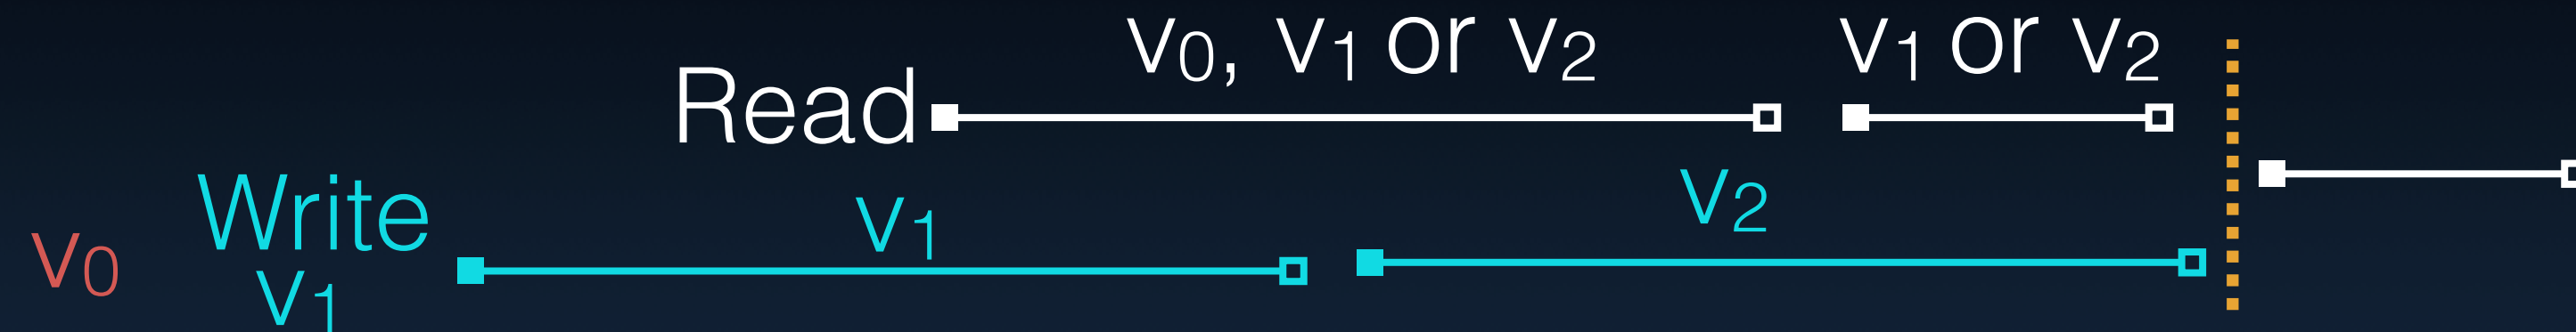
1. Each thread's history is retained in S
2. Non 'overlapping' operations retain order

Composable

Wait free

Registers

- Regular



→ Single writer

Quiescently consistent

→ Reads overlapping with a write return the newly written value or the one prior

- Atomic

A later read cannot see an earlier write

→ Multiple writers

Linearizable

→ Read the value of latest write (*takes effect* at some point in its duration)



Register Construction

Boolean \rightarrow m-valued



SW Atomic \rightarrow MW Atomic



Writers: Write in their lane,
updating clock to latest+1

Readers: Read the value with latest clock

Reg \rightarrow Atomic

(Value, Clock)

Writer: Increment Clock

Reader: Cache last read

Reuse, if new value is stale



MRSW m-valued Atomic Registers



Writers: Write $\langle \text{value}, \text{clock}++ \rangle$
Snapshot: Read all values twice
Same timestamps \Rightarrow done!

Obstruction Free

Register Snapshot

Wait free?

Writers: Make Snapshot
Write $\langle \text{value}, \text{clock}++, \text{snapshot} \rangle$

Snapshot:

Copy1 = Sequential Copy of all Register-sets
Set $\text{wrote}_i = \text{false}$ for all Registers i
Repeat until done:
 Copy2 = Sequential Copy of all Register-sets
 Sequentially compare times of Copy1 and Copy2:
 if($\text{Copy1.clock}_i \neq \text{Copy2.clock}_i$)
 if wrote_i : return Copy2.Snapshot;
 else: $\text{wrote}_i = \text{true}$; Copy1 = Copy2
return Copy1