


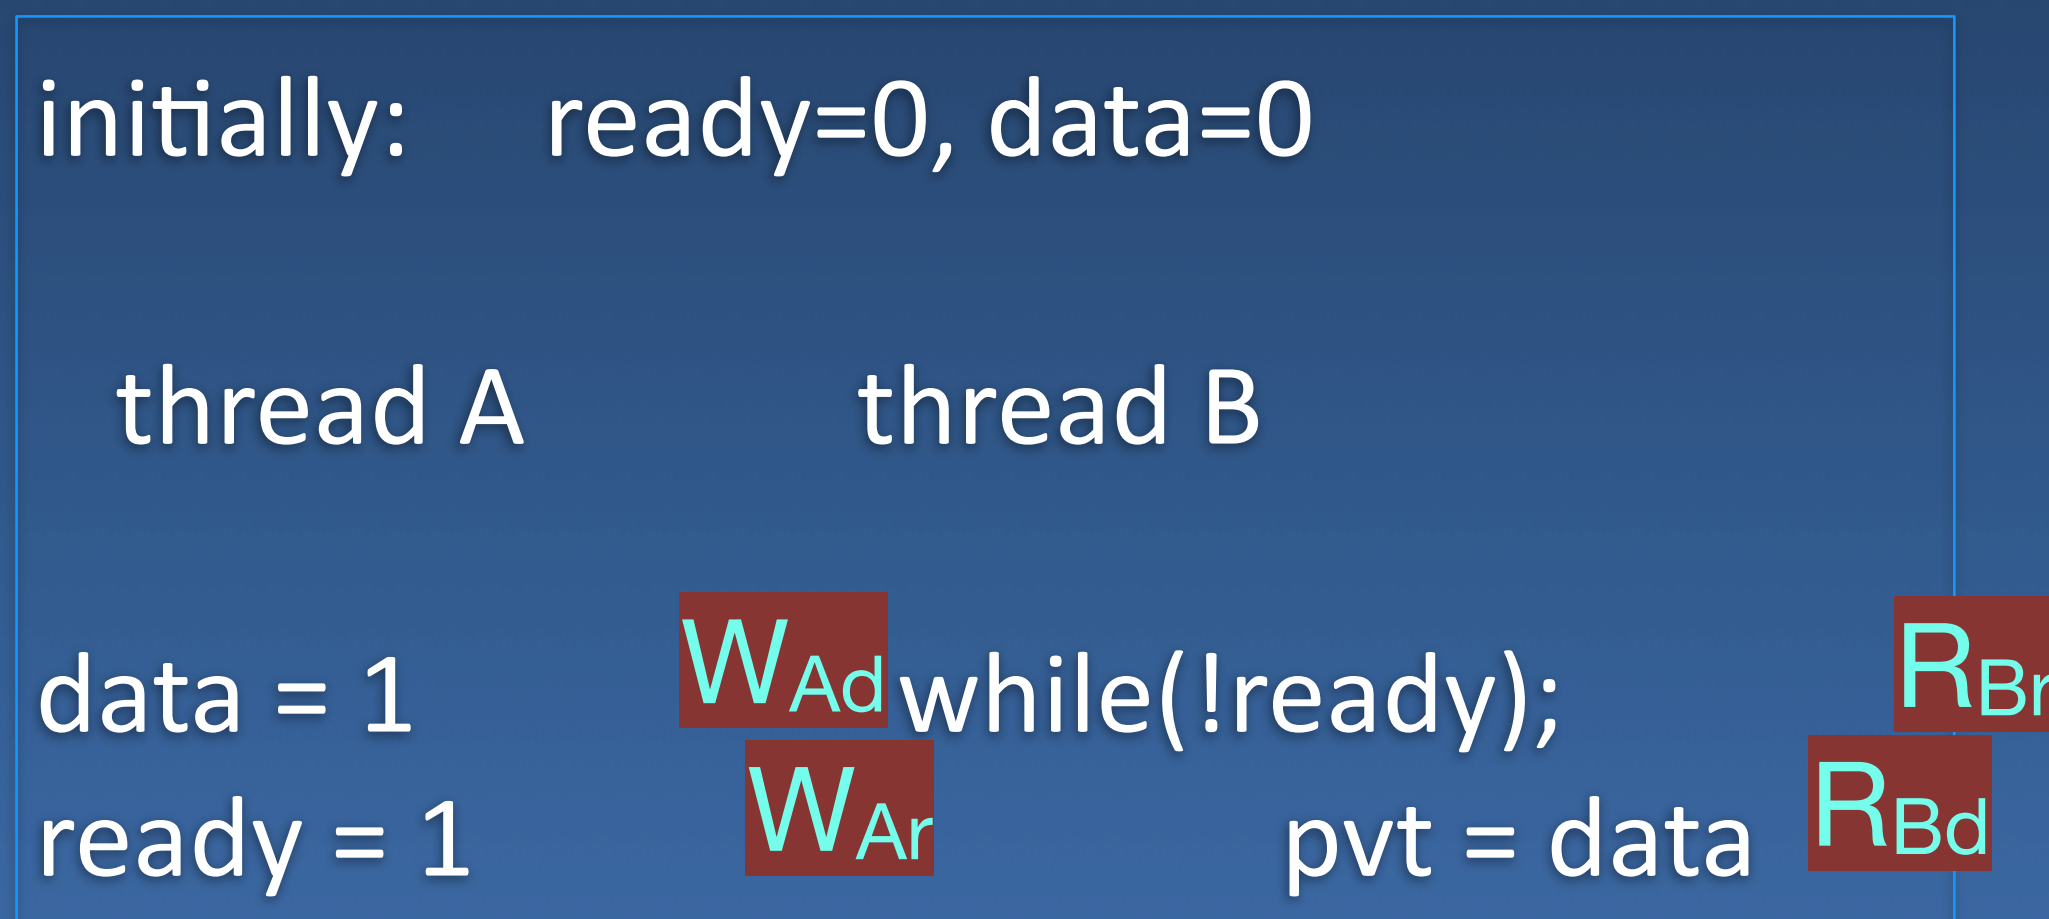
COL380

Introduction to  
Parallel & Distributed Programming

- “A multiprocessor is **sequentially consistent** if the result of any execution is the same as if the operations of all the processors were executed in **some sequential order**, and the operations of each individual processor appear in this sequence **in the order specified by its program**.” [Lamport, 1979]

# Memory Consistency Semantics


- Threads always see values written by some thread
  - ➔ No garbage
- The value seen is constrained by thread-order
  - ➔ for every thread



If B reads ready =	Then B may read data =
0	1
0	0
1	1

If B sees the new value of ready (1), it must also see the new value of data (1)

Not SC



thread A

A1 x = a

thread B

B1 x = b

thread C

C1 y1 = x (b)

C2 y2 = x (a)

thread D

D1 z1=x (a)

D2 z2=x (b)

**NOT SC**



- Hard to implement

- ➔ Poor performance

- ➔ A new memory operation cannot be issued until the previous one is completed

- ➔ No re-ordering of instructions allowed

<u>thread A</u>	<u>thread B</u>	<u>thread C</u>	<u>thread D</u>
<b>x = a</b>	<b>x = b</b>	y1 = x ( <b>b</b> )	z1=x ( <b>a</b> )
		y2 = x ( <b>a</b> )	z2=x ( <b>b</b> )

- Alternatively, out-of-order execution is allowed

- ➔ Detect and recover from SC violation

- More popular to simply switch to weaker models

- ➔ with better performance

- Hard to implement
  - ➔ Poor performance
  - ➔ A new memory operation cannot be issued until the previous one is completed
  - ➔ No re-ordering of instructions allowed
- Alternatively, out-of-order execution is allowed
  - ➔ Detect and recover from SC violation
- More popular to simply switch to weaker models
  - ➔ with better performance

# Relaxing SC

initially: ready=0, data=0

thread1

data1 = 1  
data2 = 1  
SYNC

thread 2

SYNC  
sav1 = data1  
sav2 = data2

# Causal Consistency

- Write is causally ordered after all earlier reads/writes in its thread
  - ➔ write depends on the current 'state'
- Read is causally ordered after causative write to the same variable
- Causality is transitive
- $\exists$  sequential order of causally related operations consistent with every thread's view
  - ➔ Non-related writes may be seen in different order by different threads



# Processor Consistency

- All threads see all memory writes by each thread in the order of that thread
  - ➔ all instances of write(x) are seen everywhere in the system **FIFO consistency** relaxes this constraint
  - ➔ Writes to different variables from different threads do not need to be seen in a consistent order
- Easy to implement
  - ➔ Two or more writes from a single source must remain in order, as in a pipeline
  - ➔ All writes are through to the memory

**FIFO consistency** is also known as **PRAM consistency**

## Weak Consistency

- Notion of special synchronization accesses
- Synchronization accesses are sequentially consistent
- Other accesses are consistent with synchronization accesses
  - ➔ Before any regular read/write is allowed to be visible to any other thread, all previous synchronization accesses must be completed
  - ➔ Before a synchronization access is allowed to complete, all previous ordinary read/write accesses must be completed
- Non-synchronization accesses may be re-ordered
  - ➔ suitable for optimization

- Special memory ‘fence’ operations (flush)
- Order only fences with respect to each other
- Two flushes of the same variable are synchronization operations
- Weaker than “Weak”
  - ➔ Non-intersecting flush-sets are not ordered with respect to each other



# Consistency Summary

Model	Description
Strict	Global time based atomic ordering of <i>all</i> shared accesses
Sequential	<i>All</i> threads see all shared accesses in the same order consistent with program order -- no centralized ordering
Causal	All threads see causally-related shared accesses in the same order
Processor	All threads see writes from each other in the order they were made. Writes to a variable must be seen in the same order by all threads
Weak	Special synchronization based reordering -- shared data consistent only after synchronization





What's the strictest memory consistency model being supported? Why?

