

COL380

Introduction to
Parallel & Distributed Programming

Send/Recv Synchronization

- **Synchronized**
 - ➔ Send waits for Receive to complete?
 - ➔ Synchronization (up to network delay)
- **Asynchronous**
 - ➔ Sender may proceed immediately
 - ▶ May or may not need to wait until message is copied out
 - ▶ Receiver may also proceed immediately (message arrives later)

Eager vs Rendezvous

Eager

Send-stub packetizes and transmits
(May make local copy first)
Send-stub signals Done
Recv-stub continuously accepts
Delivered when Recv call matched

Rendezvous

Send-stub transmits envelope info
May make local copy
Recv-stub continuously accepts envelope info
Recv-stub may signal OK (if it has space)
Or, wait for matching Recv call
Recv-stub sets up RDMA with Send-stub
Data transmitted
Recv-stub signals Done
Send-stub signals Done

Send Semantics

- **Standard mode:** `MPI_Send`
 - ➔ implementation dependent
 - **Buffered mode** `MPI_Bsend`
 - ➔ MPI saves a copy of message, Receiver can post later
 - ➔ User provided buffer `See MPI_Buffer_attach`
 - **Synchronous mode** `MPI_Ssend`
 - ➔ Will complete only once a matching receive has started
 - **Ready mode** `MPI_Rsend`
 - ➔ Send may start only if a matching receive has already been called
 - ➔ Helps performance
- `MPI_Send/MPI_Recv` are blocking
 - ➔ Recv blocks until output buffer is filled
 - ➔ Send blocks until some 'progress'

- In order (per pair and tag)
 - ➔ Multi-threaded applications need to be coordinate
- Progress
 - ➔ For a matching send/Recv pair, at least one of these two will complete
- Fairness not guaranteed
 - ➔ A Send or a Recv may starve because all matches are satisfied by others
- Resource limitation can cause deadlocks
- Ready/Synchronous sends requires the least resources
 - ➔ Also used for debugging

Example

If (rank == 0)

Send(to 1);

Recv(from 1);

else

Send(to 0);

Recv(from 0);

Non-blocking

- **MPI_Isend() / MPI_Irecv()**
 - ➔ Non-blocking: Control returns after setup
 - ➔ Blocking and non-blocking Send/Recv match
 - ➔ Still lower Send overhead if Recv has been posted
- **All four modes are applicable**
 - ➔ Limited impact for buffered and ready modes
- **Syntax is similar to Send and Recv**
 - ➔ MPI_Request* parameter is added to Isend and replaces MPI_Status* for IRecv

Non-blocking Call

```
int MPI_Isend(void* buf, int count, MPI_Datatype datatype, int dest, int tag,  
             MPI_Comm comm, MPI_Request *request)
```

```
int MPI_Irecv(void* buf, int count, MPI_Datatype datatype, int source, int  
             tag, MPI_Comm comm, MPI_Request *request)
```

Non-blocking calls

- **MPI_Wait(&request, &status)**

- ➔ status similar to Recv
- ➔ Blocks as per the blocking version's semantics
 - ▶ Send: message was copied out, Recv was started, etc.
 - ▶ Recv: Wait for data to fill
- ➔ Request is freed as a side-effect

Also see:

MPI_Waitany, MPI_Waitall, MPI_Waitsome
MPI_Testany, MPI_Testall, MPI_Testsome

- **MPI_Test(&request, &flag, &status)**

- ➔ Non-blocking poll
- ➔ flag indicates whether operation is complete
- ➔ Request is freed as a side-effect

Use MPI_Request_get_status to retain request

Later MPI_Request_free

- `MPI_Probe(source, tag, comm, &flag, &status)`
- `MPI_Iprobe(source, tag, comm, &flag, &status)`
 - ➔ Check information about incoming messages without actually receiving them
 - ➔ Check message size, e.g.
 - ➔ Next (matching) `Recv` will receive it
- `MPI_Cancel(&request)`
 - ➔ Request cancellation of a non-blocking request (no de-allocation)
 - ➔ Itself non-blocking: marks for cancellation and returns
 - ➔ Best effort, book-keeping necessary, can later check if cancelled

- `MPI_Send_init(buf, count, datatype, dest, tag, comm, &request);`
- `MPI_Start(&request);`
- `MPI_Start` is non-blocking
 - ➔ blocking versions do not exist
- There is also `MPI_Start_all`
 - ➔ And `MPI_Recv_init`
 - ➔ And `MPI_Bsend_init` etc.
- Reduces Process interaction with the Communication system

- **Send - Recv is point-to-point**
 - ➔ Call-to-call matching
 - ➔ Integer tag to distinguish message streams
 - ➔ Wildcard matching: MPI_ANY_SOURCE and MPI_ANY_TAG
- **Recv buffer must contain enough space for message**
 - ➔ Receiving fails otherwise
 - ➔ Can query the actual count received (MPI_Get_count)
 - ▶ Send determines the actual number sent
 - ➔ type parameter determines data<->buffer copying

MPI Data types

• MPI_CHAR	signed char
• MPI_SHORT	signed short int
• MPI_INT	signed int
• MPI_LONG	signed long int
• MPI_LONG_LONG_INT	signed long long int
• MPI_LONG_LONG	signed long long int
• MPI_SIGNED_CHAR	signed char
• MPI_UNSIGNED_CHAR	unsigned char
• MPI_UNSIGNED_SHORT	unsigned short int
• MPI_UNSIGNED	unsigned int
• MPI_UNSIGNED_LONG	unsigned long int
• MPI_UNSIGNED_LONG_LONG	unsigned long long int
• MPI_FLOAT	float
• MPI_DOUBLE	double
• MPI_LONG_DOUBLE	long double
• MPI_WCHAR	wchar_t

Objects of type
MPI_Datatype