

COL380

Introduction to  
Parallel & Distributed Programming

- MPI does not understand language's layout (struct, e.g.)

- ➔ Too system architecture dependent

MPI\_INT, MPI\_FLOAT ..

- Typemap:

- ➔ (type\_0, disp\_0), ..., (type\_n, disp\_n)

- ➔  $i^{th}$  entry is of type  $i$  and starts at byte base + disp\_ $i$

- MPI does not understand language's layout (struct, e.g.)

- ➔ Too system architecture dependent

MPI\_INT, MPI\_FLOAT ..

- Typemap:

- ➔ (type\_0, disp\_0), ..., (type\_n, disp\_n)

- ➔  $i^{th}$  entry is of type  $i$  and starts at byte base + disp\_ $i$

```
MPI_Datatype newtype;
```

```
MPI_Type_contiguous(count, MPI_INT, &newtype);
```

## Blocks

- Equally-spaced blocks of the known datatype

→ `MPI_Type_vector(blockcount, blocklength, blockstride, knowntype, &newtype);`

- ▶ Assume contiguous copies of 'knowntype'
- ▶ Stride between blocks specified in units of knowntype
- ▶ All picked blocks are of the same length



→ `MPI_Type_create_hvector(blk_count, blk_length, bytestride, knowntype, &newtype);`  
Gap between blocks is in bytes



## Blocks

- Equally-spaced blocks of the known datatype

→ `MPI_Type_vector(3blockcount, 2blocklength, 4blockstride, knowntype, &newtype);`

- ▶ Assume contiguous copies of 'knowntype'
- ▶ Stride between blocks specified in units of knowntype
- ▶ All picked blocks are of the same length



→ `MPI_Type_create_hvector(blk_count, blk_length, bytestride, knowntype, &newtype);`  
Gap between blocks is in bytes

# Generalized Blocks

- `MPI_Type_indexed(count, array_of_blocklengths, array_of_strides, knowntype, &newtype);`

- Blocks can contain different number of copies
- And may have different strides
- But the same data type



## Struct

- `MPI_Type_create_struct(count, array_of_blocklengths, array_of_bytedisplacements, array_of_knowntypes, &newtype)`

→ Example:

- ▶ Suppose `Type0 = {(double, 0), (char, 8)}`,
- ▶ `int BL[] = {2, 1, 3}, Disp[] = {0, 16, 26};`
- ▶ `MPI_Datatype Typ[] = {MPI_FLOAT, Type0, MPI_CHAR}`

→ `MPI_Type_create_struct(3, BL, Disp, Typ, &newtype):`

- ▶ `(float, 0), (float, 4), (double, 16), (char, 24), (char, 26), (char, 27), (char, 28)`






- `MPI_Type_create_struct(count, array_of_blocklengths, array_of_bytedisplacements, array_of_knowntypes, &newtype)`

→ Example:

▶ Suppose `Type0 = {(double, 0), (char, 8)}`,



▶ `int BL[] = {2, 1, 3}, Disp[] = {0, 16, 26};`

▶ `MPI_Datatype Typ[] = {MPI_FLOAT, Type0, MPI_CHAR}`

→ `MPI_Type_create_struct(3, BL, Disp, Typ, &newtype):`


▶ `(float, 0), (float, 4), (double, 16), (char, 24), (char, 26), (char, 27), (char, 28)`



- `MPI_Type_create_struct(count, array_of_blocklengths, array_of_bytedisplacements, array_of_knowntypes, &newtype)`

→ Example:

▶ Suppose `Type0 = {(double, 0), (char, 8)}`,



▶ `int BL[] = {2, 1, 3}, Disp[] = {0, 16, 26};`



▶ `MPI_Datatype Typ[] = {MPI_FLOAT, Type0, MPI_CHAR}`


→ `MPI_Type_create_struct(3, BL, Disp, Typ, &newtype):`

▶ `(float, 0), (float, 4), (double, 16), (char, 24), (char, 26), (char, 27), (char, 28)`


- `MPI_Type_create_struct(count, array_of_blocklengths, array_of_bytedisplacements, array_of_knowntypes, &newtype)`

→ Example:

▶ Suppose `Type0 = {(double, 0), (char, 8)}`,



▶ `int BL[] = {2, 1, 3}, Disp[] = {0, 16, 26};`



▶ `MPI_Datatype Typ[] = {MPI_FLOAT, Type0, MPI_CHAR}`


→ `MPI_Type_create_struct(3, BL, Disp, Typ, &newtype):`

▶ `(float, 0), (float, 4), (double, 16), (char, 24), (char, 26), (char, 27), (char, 28)`

- `MPI_Type_create_struct(count, array_of_blocklengths, array_of_bytedisplacements, array_of_knowntypes, &newtype)`

→ Example:

▶ Suppose `Type0 = {(double, 0), (char, 8)}`,



▶ `int BL[] = {2, 1, 3}, Disp[] = {0, 16, 26};`



▶ `MPI_Datatype Typ[] = {MPI_FLOAT, Type0, MPI_CHAR}`

→ `MPI_Type_create_struct(3, BL, Disp, Typ, &newtype):`


▶ `(float, 0), (float, 4), (double, 16), (char, 24), (char, 26), (char, 27), (char, 28)`



- `MPI_Type_create_struct(count, array_of_blocklengths, array_of_bytedisplacements, array_of_knowntypes, &newtype)`

→ Example:

▶ Suppose `Type0 = {(double, 0), (char, 8)}`,



▶ `int BL[] = {2, 1, 3}, Disp[] = {0, 16, 26};`



▶ `MPI_Datatype Typ[] = {MPI_FLOAT, Type0, MPI_CHAR}`

→ `MPI_Type_create_struct(3, BL, Disp, Typ, &newtype):`

▶ `(float, 0), (float, 4), (double, 16), (char, 24), (char, 26), (char, 27), (char, 28)`

`MPI_Type_get_contents(..)`



- `MPI_Type_commit(&datatype)`
  - ➔ A datatype object must be committed before communication
- `MPI_Type_size(datatype, &size)`
  - ➔ Total size in bytes
- `MPI_Type_get_extent(datatype, &beg, &extent);`
- `MPI_Type_create_resized(datatype, beg, extent, &newtype);`
- `MPI_Get_address(data, &Address[0]);`
- `MPI_BOTTOM`

## Data Type Functions

- `MPI_Type_commit(&datatype)`

→ A datatype object must be committed before communication

- `MPI_Type_size(datatype, &size)`

→ Total size in bytes

- `MPI_Type_get_extent(datatype, &b`

```
MPI_Datatype atype;  
MPI_Type_contiguous(4, MPI_CHAR, &atype);  
int asize;  
MPI_Type_size(atype, &asize);  
MPI_Type_commit(&atype);  
MPI_Send(buf, nItems, atype, dest, ..);  
MPI_Recv(...);
```

- `MPI_Type_create_resized(datatype, beg, extent, &newtype);`

- `MPI_Get_address(data, &Address[0]);`

- `MPI_BOTTOM`

- `MPI_Type_commit(&datatype)`
  - ➔ A datatype object must be committed before communication
- `MPI_Type_size(datatype, &size)`
  - ➔ Total size in bytes
- `MPI_Type_get_extent(datatype, &beg, &extent);`
- `MPI_Type_create_resized(datatype, beg, extent, &newtype);`
- `MPI_Get_address(data, &Address[0]);`
- `MPI_BOTTOM`

# Derived Datatype

```
struct Particle
{
    int class;    // particle class
    double d[6]; // particle coordinates
    char b[7];    // some additional info
};
```



## Derived Datatype

```
sendParticles(struct Particle particle[], int N):
```

```
    MPI_Datatype Particletype;
```

```
    MPI_Datatype types[3] = {MPI_INT, MPI_DOUBLE, MPI_CHAR};
```

```
    int blockcount[3] = {1, 6, 7};
```

```
    /* compute displacements of structure components */
```

```
    MPI_Aint disp[3];
```

```
    MPI_Address(particle, disp);
```

```
    MPI_Address(particle[0].d, disp+1);
```

```
    MPI_Address(particle[0].b, disp+2);
```

```
    for (int i=2; i >= 0; i--) disp[i] -= disp[0];
```

```
    MPI_Type_struct(3, blockcount, disp, types, &Particletype);
```

```
    MPI_Type_commit( &Particletype);
```

```
    MPI_Send(particle, N, Particletype, dest, tag, comm);
```

```
struct Particle
```

```
{
```


```
    int class;    // particle class
```

```
    double d[6]; // particle coordinates
```



```
    char b[7];    // some additional info
```

```
};
```


# Data Transfer



# Data Transfer



## Data Transfer



```
MPI_Status status;
MPI_Datatype column;
MPI_Type_vector(n, 1, n+1, MPI_INT, &column);
MPI_Type_commit(&column);
if(rank == 0) {
    MPI_Send(A+n-1, 1, column, 1, tag, MPI_COMM_WORLD);
    MPI_Recv(A+n, 1, column, 1, tag, MPI_COMM_WORLD, &status);
}
if(rank == 1) {
    MPI_Recv(A, 1, column, 0, tag, MPI_COMM_WORLD, &status);
    MPI_Send(A+1, 1, column, 0, tag, MPI_COMM_WORLD);
}
```

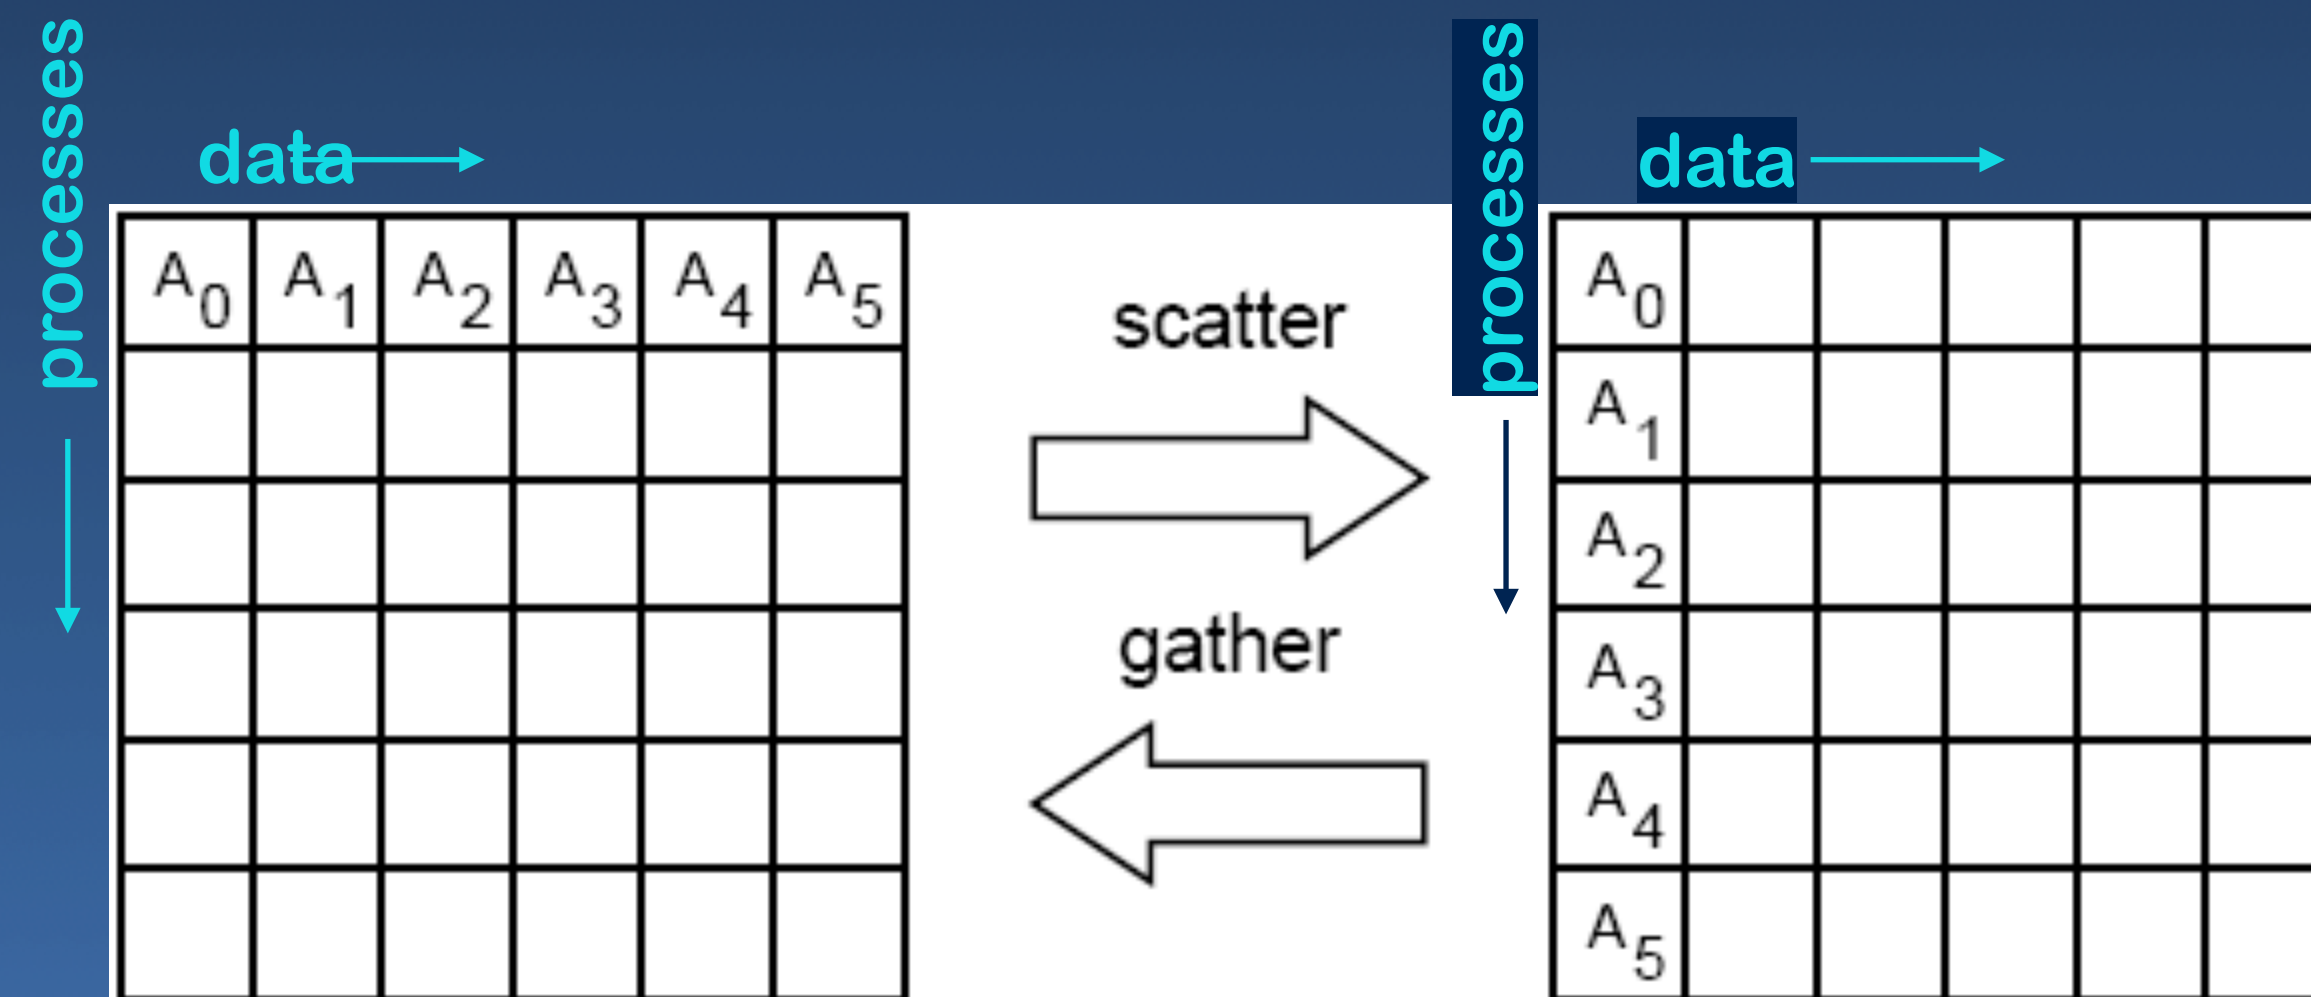
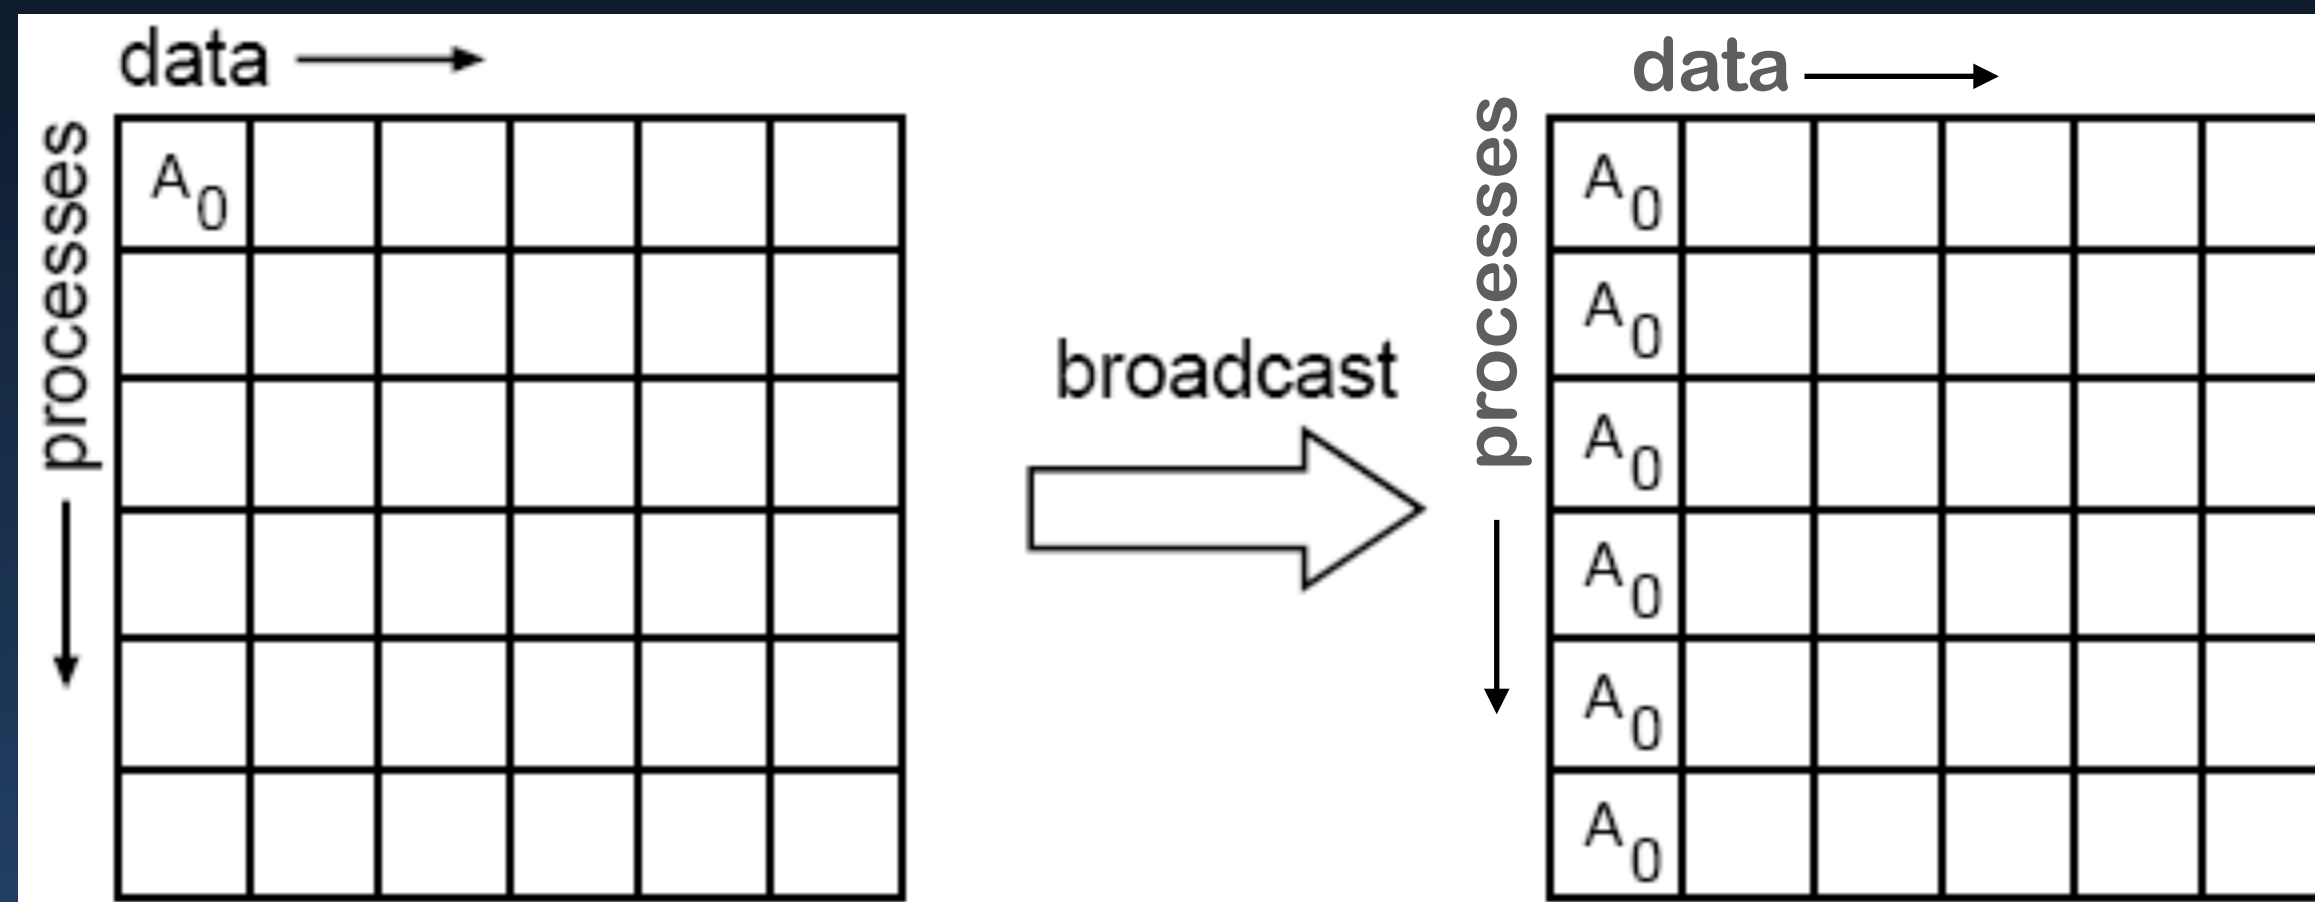


- MPI\_Barrier
  - Barrier synchronization across all members of a group
- MPI\_Bcast
  - Broadcast from one member to all members of a group
- MPI\_Scatter, MPI\_Gather, MPI\_Allgather
  - Gather data from all members of a group to one
- MPI\_Alltoall
  - complete exchange or all-to-all
- MPI\_Allreduce, MPI\_Reduce
  - Reduction operations
- MPI\_Reduce\_Scatter
  - Combined reduction and scatter operation
- MPI\_Scan, MPI\_Exscan
  - Prefix

- Synchronization of the calling processes
  - the call blocks until all of the processes have placed the call


```
MPI_Barrier(comm) ;
```

# Collective Communication



Example from MPI document

# Collective Communication





## Broadcast

- Broadcast: one sender, many receivers
- Includes all processes in communicator
  - ➔ all processes must make a call to MPI\_Bcast
  - ➔ Must agree on sender
- Broadcast does not mandate global synchronization
  - ➔ Some implementations may incur synchronization
  - ➔ Call may return before other have received, e.g.
  - ➔ Different from MPI\_Barrier(communicator)

## Broadcast

```
MPI_Bcast(msg, count, MPI_INT, root, comm);
```

**msg**        pointer to message buffer

**count**     number of items sent

**MPI\_INT**   type of item sent

**root**       sending processor

- **Again:** All participants must call
- count and type should be the same on all members
- Can broadcast on inter-communicators also

# Deadlock?

- Thread 0:  
    MPI\_Bcast(buf1, count, type, 0, comm);  
    MPI\_Bcast(buf2, count, type, 1, comm);
- Thread 1:  
    MPI\_Bcast(buf1, count, type, 1, comm);  
    MPI\_Bcast(buf2, count, type, 0, comm);

# Deadlock?

- Thread 0:

```
MPI_Bcast(buf1, count, type, 0, comm);  
MPI_Bcast(buf2, count, type, 1, comm);
```

- Thread 1:

```
MPI_Bcast(buf1, count, type, 1, comm);  
MPI_Bcast(buf2, count, type, 0, comm);
```





# Deadlock?

- Thread 0:

`MPI_Bcast(buf1, count, type, 0, comm);`

`MPI_Bcast(buf2, count, type, 1, comm);`

- Thread 1:

`MPI_Bcast(buf1, count, type, 1, comm);`

`MPI_Bcast(buf2, count, type, 0, comm);`



```
MPI_Gather(sendbuf, sendcount, sendtype,  
recvbuf, recvcount, recvtype, root, comm);
```

- Similar to non-roots sending:
  - `MPI_Send(sendbuf, sendcount, sendtype, root, ...)`,
- and the root receiving n times:
  - `MPI_Recv(recvbuf + i * recvcount * extent(recvtype), recvcount, recvtype, i, ...)`,
- **MPI\_Gatherv** allows different size data to be gathered
- **MPI\_Allgather** has No root, all nodes get result

## Gather Example


```
MPI_Comm com;  
int gsize, sendarray[100];  
int root, *recvbuf;  
MPI_Datatype rtype;  
...  
MPI_Comm_size( comm, &gsize);  
MPI_Type_contiguous( 100, MPI_INT, &rtype );  
MPI_Type_commit( &rtype );  
recvbuf = (int *) malloc(gsize * 100 * sizeof(int));  
MPI_Gather(sendarray, 100, MPI_INT, recvbuf, 1, rtype, root, comm)
```



```
double A[8][8], alocal[4][4];
int i, j, r, rank, size, sendcount[4], sdispls[4];
MPI_Datatype stype, vtype;

MPI_Comm_rank( MPI_COMM_WORLD, &rank );
MPI_Comm_size( MPI_COMM_WORLD, &size );
if (size != 4) MPI_Abort( MPI_COMM_WORLD, 1 );
if (rank == 0) {
```

## Scatter Matrix



```
    initialize(A);
    MPI_Type_vector(4, 4, 8, MPI_DOUBLE, &vtype ); // 4 sets of 4 doubles, separated by 8
    MPI_Type_create_resized(vtype, 0, 4*sizeof(double), &stype); // Artificial type for scatter
    MPI_Type_commit(&stype );

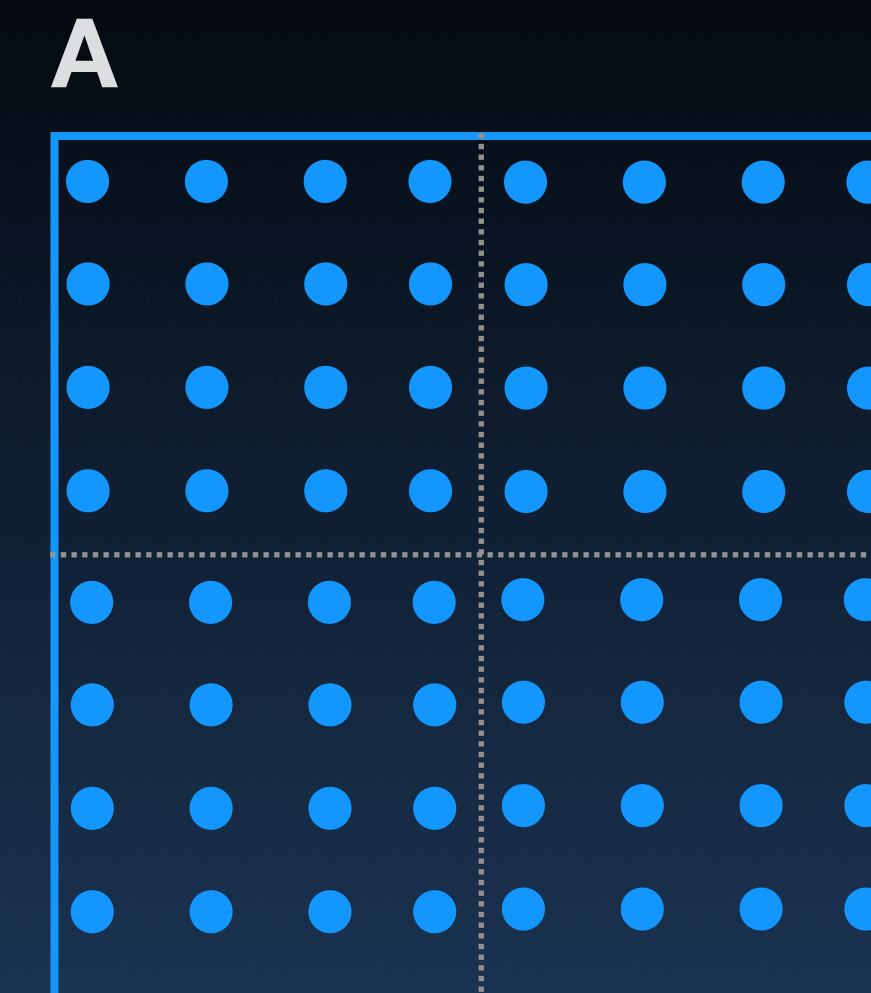
    // Setup the Scatter values for the send buffer
    sendcount[0] = sendcount[1] = sendcount[2] = sendcount[3] = 1; // Send one to each
    // Starting locations in A of the four sub matrices in terms of stype
    sdispls[0] = 0; sdispls[1] = 1; sdispls[2] = 8; sdispls[3] = 9;
    MPI_Scatterv(A, sendcount, sdispls, stype, alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
} else {
    MPI_Scatterv( (void *)0, (void *)0, (void *)0, MPI_DATATYPE_NULL,
                  alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
}
```



```
double A[8][8], alocal[4][4];
int i, j, r, rank, size, sendcount[4], sdispls[4];
MPI_Datatype stype, vtype;

MPI_Comm_rank( MPI_COMM_WORLD, &rank );
MPI_Comm_size( MPI_COMM_WORLD, &size );
if (size != 4) MPI_Abort( MPI_COMM_WORLD, 1 );
if (rank == 0) {
```

## Scatter Matrix




```
    initialize(A);
    MPI_Type_vector(4, 4, 8, MPI_DOUBLE, &vtype ); // 4 sets of 4 doubles, separated by 8
    MPI_Type_create_resized(vtype, 0, 4*sizeof(double), &stype); // Artificial type for scatter
    MPI_Type_commit(&stype );

    // Setup the Scatter values for the send buffer
    sendcount[0] = sendcount[1] = sendcount[2] = sendcount[3] = 1; // Send one to each
    // Starting locations in A of the four sub matrices in terms of stype
    sdispls[0] = 0; sdispls[1] = 1; sdispls[2] = 8; sdispls[3] = 9;
    MPI_Scatterv(A, sendcount, sdispls, stype, alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
} else {
    MPI_Scatterv( (void *)0, (void *)0, (void *)0, MPI_DATATYPE_NULL,
                  alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
}
```

```
double A[8][8], alocal[4][4];
int i, j, r, rank, size, sendcount[4], sdispls[4];
MPI_Datatype stype, vtype;

MPI_Comm_rank( MPI_COMM_WORLD, &rank );
MPI_Comm_size( MPI_COMM_WORLD, &size );
if (size != 4) MPI_Abort( MPI_COMM_WORLD, 1 );
if (rank == 0) {
```

## Scatter Matrix



```
    initialize(A);
    MPI_Type_vector(4, 4, 8, MPI_DOUBLE, &vtype ); // 4 sets of 4 doubles, separated by 8
    MPI_Type_create_resized(vtype, 0, 4*sizeof(double), &stype); // Artificial type for scatter
    MPI_Type_commit(&stype );


    // Setup the Scatter values for the send buffer
    sendcount[0] = sendcount[1] = sendcount[2] = sendcount[3] = 1; // Send one to each
    // Starting locations in A of the four sub matrices in terms of stype
    sdispls[0] = 0; sdispls[1] = 1; sdispls[2] = 8; sdispls[3] = 9;
    MPI_Scatterv(A, sendcount, sdispls, stype, alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
} else {
    MPI_Scatterv( (void *)0, (void *)0, (void *)0, MPI_DATATYPE_NULL,
                  alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
}
```



```
double A[8][8], alocal[4][4];
int i, j, r, rank, size, sendcount[4], sdispls[4];
MPI_Datatype stype, vtype;

MPI_Comm_rank( MPI_COMM_WORLD, &rank );
MPI_Comm_size( MPI_COMM_WORLD, &size );
if (size != 4) MPI_Abort( MPI_COMM_WORLD, 1 );
if (rank == 0) {
```

## Scatter Matrix




```
    initialize(A);
    MPI_Type_vector(4, 4, 8, MPI_DOUBLE, &vtype ); // 4 sets of 4 doubles, separated by 8
    MPI_Type_create_resized(vtype, 0, 4*sizeof(double), &stype); // Artificial type for scatter
    MPI_Type_commit(&stype );

    // Setup the Scatter values for the send buffer
    sendcount[0] = sendcount[1] = sendcount[2] = sendcount[3] = 1; // Send one to each
    // Starting locations in A of the four sub matrices in terms of stype
    sdispls[0] = 0; sdispls[1] = 1; sdispls[2] = 8; sdispls[3] = 9;
    MPI_Scatterv(A, sendcount, sdispls, stype, alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
} else {
    MPI_Scatterv( (void *)0, (void *)0, (void *)0, MPI_DATATYPE_NULL,
                  alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
}
```

```
double A[8][8], alocal[4][4];
int i, j, r, rank, size, sendcount[4], sdispls[4];
MPI_Datatype stype, vtype;

MPI_Comm_rank( MPI_COMM_WORLD, &rank );
MPI_Comm_size( MPI_COMM_WORLD, &size );
if (size != 4) MPI_Abort( MPI_COMM_WORLD, 1 );
if (rank == 0) {
```

## Scatter Matrix



```
    initialize(A);
    MPI_Type_vector(4, 4, 8, MPI_DOUBLE, &vtype ); // 4 sets of 4 doubles, separated by 8
    MPI_Type_create_resized(vtype, 0, 4*sizeof(double), &stype); // Artificial type for scatter
    MPI_Type_commit(&stype );


    // Setup the Scatter values for the send buffer
    sendcount[0] = sendcount[1] = sendcount[2] = sendcount[3] = 1; // Send one to each
    // Starting locations in A of the four sub matrices in terms of stype
    sdispls[0] = 0; sdispls[1] = 1; sdispls[2] = 8; sdispls[3] = 9;
    MPI_Scatterv(A, sendcount, sdispls, stype, alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
} else {
    MPI_Scatterv( (void *)0, (void *)0, (void *)0, MPI_DATATYPE_NULL,
                  alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
}
```



```
double A[8][8], alocal[4][4];
int i, j, r, rank, size, sendcount[4], sdispls[4];
MPI_Datatype stype, vtype;

MPI_Comm_rank( MPI_COMM_WORLD, &rank );
MPI_Comm_size( MPI_COMM_WORLD, &size );
if (size != 4) MPI_Abort( MPI_COMM_WORLD, 1 );
if (rank == 0) {
```

## Scatter Matrix




```
    initialize(A);
    MPI_Type_vector(4, 4, 8, MPI_DOUBLE, &vtype ); // 4 sets of 4 doubles, separated by 8
    MPI_Type_create_resized(vtype, 0, 4*sizeof(double), &stype); // Artificial type for scatter
    MPI_Type_commit(&stype );

    // Setup the Scatter values for the send buffer
    sendcount[0] = sendcount[1] = sendcount[2] = sendcount[3] = 1; // Send one to each
    // Starting locations in A of the four sub matrices in terms of stype
    sdispls[0] = 0; sdispls[1] = 1; sdispls[2] = 8; sdispls[3] = 9;
    MPI_Scatterv(A, sendcount, sdispls, stype, alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
} else {
    MPI_Scatterv( (void *)0, (void *)0, (void *)0, MPI_DATATYPE_NULL,
                  alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
}
```

```
double A[8][8], alocal[4][4];
int i, j, r, rank, size, sendcount[4], sdispls[4];
MPI_Datatype stype, vtype;

MPI_Comm_rank( MPI_COMM_WORLD, &rank );
MPI_Comm_size( MPI_COMM_WORLD, &size );
if (size != 4) MPI_Abort( MPI_COMM_WORLD, 1 );
if (rank == 0) {
```

## Scatter Matrix



```
    initialize(A);
    MPI_Type_vector(4, 4, 8, MPI_DOUBLE, &vtype ); // 4 sets of 4 doubles, separated by 8
    MPI_Type_create_resized(vtype, 0, 4*sizeof(double), &stype); // Artificial type for scatter
    MPI_Type_commit(&stype );


    // Setup the Scatter values for the send buffer
    sendcount[0] = sendcount[1] = sendcount[2] = sendcount[3] = 1; // Send one to each
    // Starting locations in A of the four sub matrices in terms of stype
    sdispls[0] = 0; sdispls[1] = 1; sdispls[2] = 8; sdispls[3] = 9;
    MPI_Scatterv(A, sendcount, sdispls, stype, alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
} else {
    MPI_Scatterv( (void *)0, (void *)0, (void *)0, MPI_DATATYPE_NULL,
                  alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
}
```



```
double A[8][8], alocal[4][4];
int i, j, r, rank, size, sendcount[4], sdispls[4];
MPI_Datatype stype, vtype;

MPI_Comm_rank( MPI_COMM_WORLD, &rank );
MPI_Comm_size( MPI_COMM_WORLD, &size );
if (size != 4) MPI_Abort( MPI_COMM_WORLD, 1 );
if (rank == 0) {
```

## Scatter Matrix




```
    initialize(A);
    MPI_Type_vector(4, 4, 8, MPI_DOUBLE, &vtype ); // 4 sets of 4 doubles, separated by 8
    MPI_Type_create_resized(vtype, 0, 4*sizeof(double), &stype); // Artificial type for scatter
    MPI_Type_commit(&stype );

    // Setup the Scatter values for the send buffer
    sendcount[0] = sendcount[1] = sendcount[2] = sendcount[3] = 1; // Send one to each
    // Starting locations in A of the four sub matrices in terms of stype
    sdispls[0] = 0; sdispls[1] = 1; sdispls[2] = 8; sdispls[3] = 9;
    MPI_Scatterv(A, sendcount, sdispls, stype, alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
} else {
    MPI_Scatterv( (void *)0, (void *)0, (void *)0, MPI_DATATYPE_NULL,
                  alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
}
```

```
double A[8][8], alocal[4][4];
int i, j, r, rank, size, sendcount[4], sdispls[4];
MPI_Datatype stype, vtype;

MPI_Comm_rank( MPI_COMM_WORLD, &rank );
MPI_Comm_size( MPI_COMM_WORLD, &size );
if (size != 4) MPI_Abort( MPI_COMM_WORLD, 1 );
if (rank == 0) {
```

## Scatter Matrix



```
    initialize(A);
    MPI_Type_vector(4, 4, 8, MPI_DOUBLE, &vtype ); // 4 sets of 4 doubles, separated by 8
    MPI_Type_create_resized(vtype, 0, 4*sizeof(double), &stype); // Artificial type for scatter
    MPI_Type_commit(&stype );


    // Setup the Scatter values for the send buffer
    sendcount[0] = sendcount[1] = sendcount[2] = sendcount[3] = 1; // Send one to each
    // Starting locations in A of the four sub matrices in terms of stype
    sdispls[0] = 0; sdispls[1] = 1; sdispls[2] = 8; sdispls[3] = 9;
    MPI_Scatterv(A, sendcount, sdispls, stype, alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
} else {
    MPI_Scatterv( (void *)0, (void *)0, (void *)0, MPI_DATATYPE_NULL,
                  alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
}
```



```
double A[8][8], alocal[4][4];
int i, j, r, rank, size, sendcount[4], sdispls[4];
MPI_Datatype stype, vtype;

MPI_Comm_rank( MPI_COMM_WORLD, &rank );
MPI_Comm_size( MPI_COMM_WORLD, &size );
if (size != 4) MPI_Abort( MPI_COMM_WORLD, 1 );
if (rank == 0) {
```

## Scatter Matrix



```
    initialize(A);
    MPI_Type_vector(4, 4, 8, MPI_DOUBLE, &vtype ); // 4 sets of 4 doubles, separated by 8
    MPI_Type_create_resized(vtype, 0, 4*sizeof(double), &stype); // Artificial type for scatter
    MPI_Type_commit(&stype );

    // Setup the Scatter values for the send buffer
    sendcount[0] = sendcount[1] = sendcount[2] = sendcount[3] = 1; // Send one to each
    // Starting locations in A of the four sub matrices in terms of stype
    sdispls[0] = 0; sdispls[1] = 1; sdispls[2] = 8; sdispls[3] = 9;
    MPI_Scatterv(A, sendcount, sdispls, stype, alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
} else {
    MPI_Scatterv( (void *)0, (void *)0, (void *)0, MPI_DATATYPE_NULL,
                  alocal, 4*4, MPI_DOUBLE, 0, MPI_COMM_WORLD );
}
```

```
MPI_Reduce(dataArray, resultArray, count,  
           type, MPI_SUM, root, com);
```

<b>dataArray</b>	data sent from each processor
<b>Result</b>	stores result of combining operation
<b>count</b>	number of items in each of dataArray, result
<b>MPI_SUM</b>	combining operation, one of a predefined set
<b>root</b>	rank of processor receiving data

- Multiple elements can be reduced in one shot
- Illegal to alias input and output arrays

- MPI\_Reduce: result is at the root
  - ➔ operation repeated for each element of the input arrays on each processor
- MPI\_Allreduce: result is sent out to everyone
- MPI\_Reduce\_scatter: equivalent to a reduce followed by a scatter
- User defined operations



## User-defined reduce operation

```
void rfunction(void *invec, void *inoutvec, int *len,  
               MPI_Datatype *datatype){  
    // accumulate *len type items of invec into inoutvec  
}
```

```
MPI_Op op;  
MPI_Op_create(rfunction, commute, &op);  
MPI_Reduce(inArray, outArray, count, type, op, root, com);
```

Later:

```
MPI_op_free(&op);
```




```
MPI_Scan(sendbuf, recvbuf, count, datatype, op, comm);
```

- Prefix reduction on data in sendbuf
  - ➔ Multiple prefix ops in one shot
- Returns in the receive buffer of the process i:
  - ➔ reduction of the values in the send buffers of processes 0,...,i (inclusive)
- All ranks must agree on op, datatype, count
- **MPI\_EScan** for exclusive scan

## In-place MPI\_Scan

```
MPI_Scan(MPI_IN_PLACE, recvbuf, 5, MPI_INT, MPI_SUM, comm);
```

	recvbuf				
P0	3	4	2	8	1
P1	5	2	5	1	7
P2	2	4	4	10	4
P3	1	6	9	3	1



	recvbuf				
P0	3	4	2	8	1
P1	8	6	7	9	8
P2	10	10	11	19	12
P3	11	16	12	22	13

- `MPI_Comm_Spawn(command, argv, maxprocs, info, root, comm, intercomm, array_of_errcodes)`
- The children have their own `MPI_COMM_WORLD`
- May not return until `MPI_INIT` has been called in the children
- More efficient to start all processes at once



```
MPI_Win_create(basemem, size, displ_unit, info, MPI_COMM_WORLD,  
&win);
```

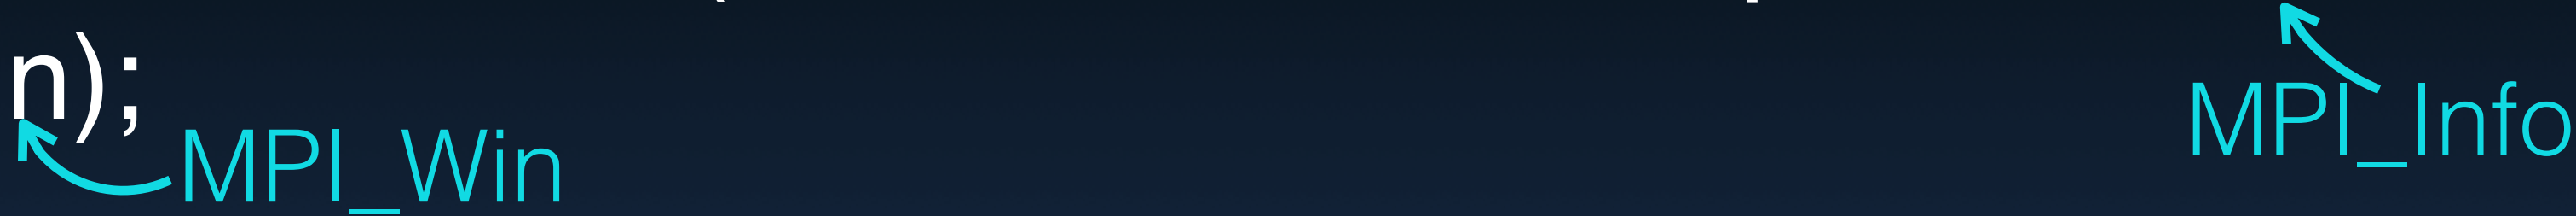


Diagram illustrating the parameters of `MPI_Win_create`:  
• `&win` is the address of the `MPI_Win` object.  
• `info` is a pointer to an `MPI_Info` object.

...

```
MPI_Win_free(&win);
```

- Weak synchronization
- Collective call
- Info specifies system-specific information (e.g., memory locking)
  - ➔ Designed for optimizing performance
- See `MPI_Alloc_mem/MPI_Win_allocate` for `basemem` allocation

- `MPI_Put(my_addr, my_count, my_datatype, there_rank, there_disp, there_count, there_datatype, win);`
  - Written in the dest window-buffer at address
    - ▶  $\text{window\_base} + \text{disp} \times \text{disp\_unit}$
  - Must fit in the target buffer
  - `there_datatype` defined on the “putter”
    - ▶ But refers to memory “there”
    - ▶ Usually defined on both sides

- `MPI_Put(my_addr, my_count, my_datatype, there_rank, there_disp, there_count, there_datatype, win);`
  - Written in the dest window-buffer at address
    - ▶  $\text{window\_base} + \text{disp} \times \text{disp\_unit}$
  - Must fit in the target buffer
  - `there_datatype` defined on the “putter”
    - ▶ But refers to memory “there”
    - ▶ Usually defined on both sides

MPI\_Get does the reverse: there → my

Also see:

`MPI_Accumulate`

performs an “op” at destination



# Remote Memory Synchronization

- MPI\_Win\_fence
- MPI\_Win\_flush
- MPI\_Win\_lock
- MPI\_Win\_unlock
- MPI\_Win\_start
- MPI\_Win\_complete
- MPI\_Win\_post
- MPI\_Win\_Wait
- MPI\_Win\_Test

## Remote Memory Synchronization

- MPI\_Win\_fence
- MPI\_Win\_flush
- MPI\_Win\_lock
- MPI\_Win\_unlock
- MPI\_Win\_start
- MPI\_Win\_complete
- MPI\_Win\_post
- MPI\_Win\_Wait
- MPI\_Win\_Test

```
int winbuf[10];
MPI_Win windo;
MPI_Win_create(winbuf, 10*sizeof(int), sizeof(int),
               MPI_INFO_NULL, MPI_COMM_WORLD, &windo);
MPI_Win_fence(0, windo); // Collective
```

```
if(rank == 1) {
    int lbuf[5];
    initialize(lbuf);
    MPI_Put(lbuf, 5, MPI_INT, 0, 5, 5, MPI_INT, windo);
}
```

```
MPI_Win_fence(0, windo); // Wait for MPI_Put complete
```

```
if(my_rank == 0)
    use(winbuf+5);
```

# Remote Memory Synchronization

- MPI\_Win\_fence
- MPI\_Win\_flush
- MPI\_Win\_lock
- MPI\_Win\_unlock
- MPI\_Win\_start
- MPI\_Win\_complete
- MPI\_Win\_post
- MPI\_Win\_Wait
- MPI\_Win\_Test

```
int winbuf[10];
MPI_Win windo;
MPI_Win_create(winbuf, 10*sizeof(int), sizeof(int),
               MPI_INFO_NULL, MPI_COMM_WORLD, &windo);
MPI_Win_fence(0, windo); // Collective
```

```
if(rank == 1) {
    int lbuf[5];
    initialize(lbuf);
    MPI_Put(lbuf, 5, MPI_INT, 0, 5, 5, MPI_INT, windo);
}
```

“Assert”



```
MPI_Win_fence(0, windo); // Wait for MPI_Put complete
```

```
if(my_rank == 0)
    use(winbuf+5);
```



# Remote Memory Synchronization

- MPI\_Win\_fence
- MPI\_Win\_flush
- MPI\_Win\_lock
- MPI\_Win\_unlock
- MPI\_Win\_start
- MPI\_Win\_complete
- MPI\_Win\_post
- MPI\_Win\_Wait
- MPI\_Win\_Test

**Look these up**

```
int winbuf[10];
MPI_Win windo;
MPI_Win_create(winbuf, 10*sizeof(int), sizeof(int),
               MPI_INFO_NULL, MPI_COMM_WORLD, &windo);
MPI_Win_fence(0, windo); // Collective
```

```
if(rank == 1) {
    int lbuf[5];
    initialize(lbuf);
    MPI_Put(lbuf, 5, MPI_INT, 0, 5, 5, MPI_INT, windo);
}
```

“Assert”

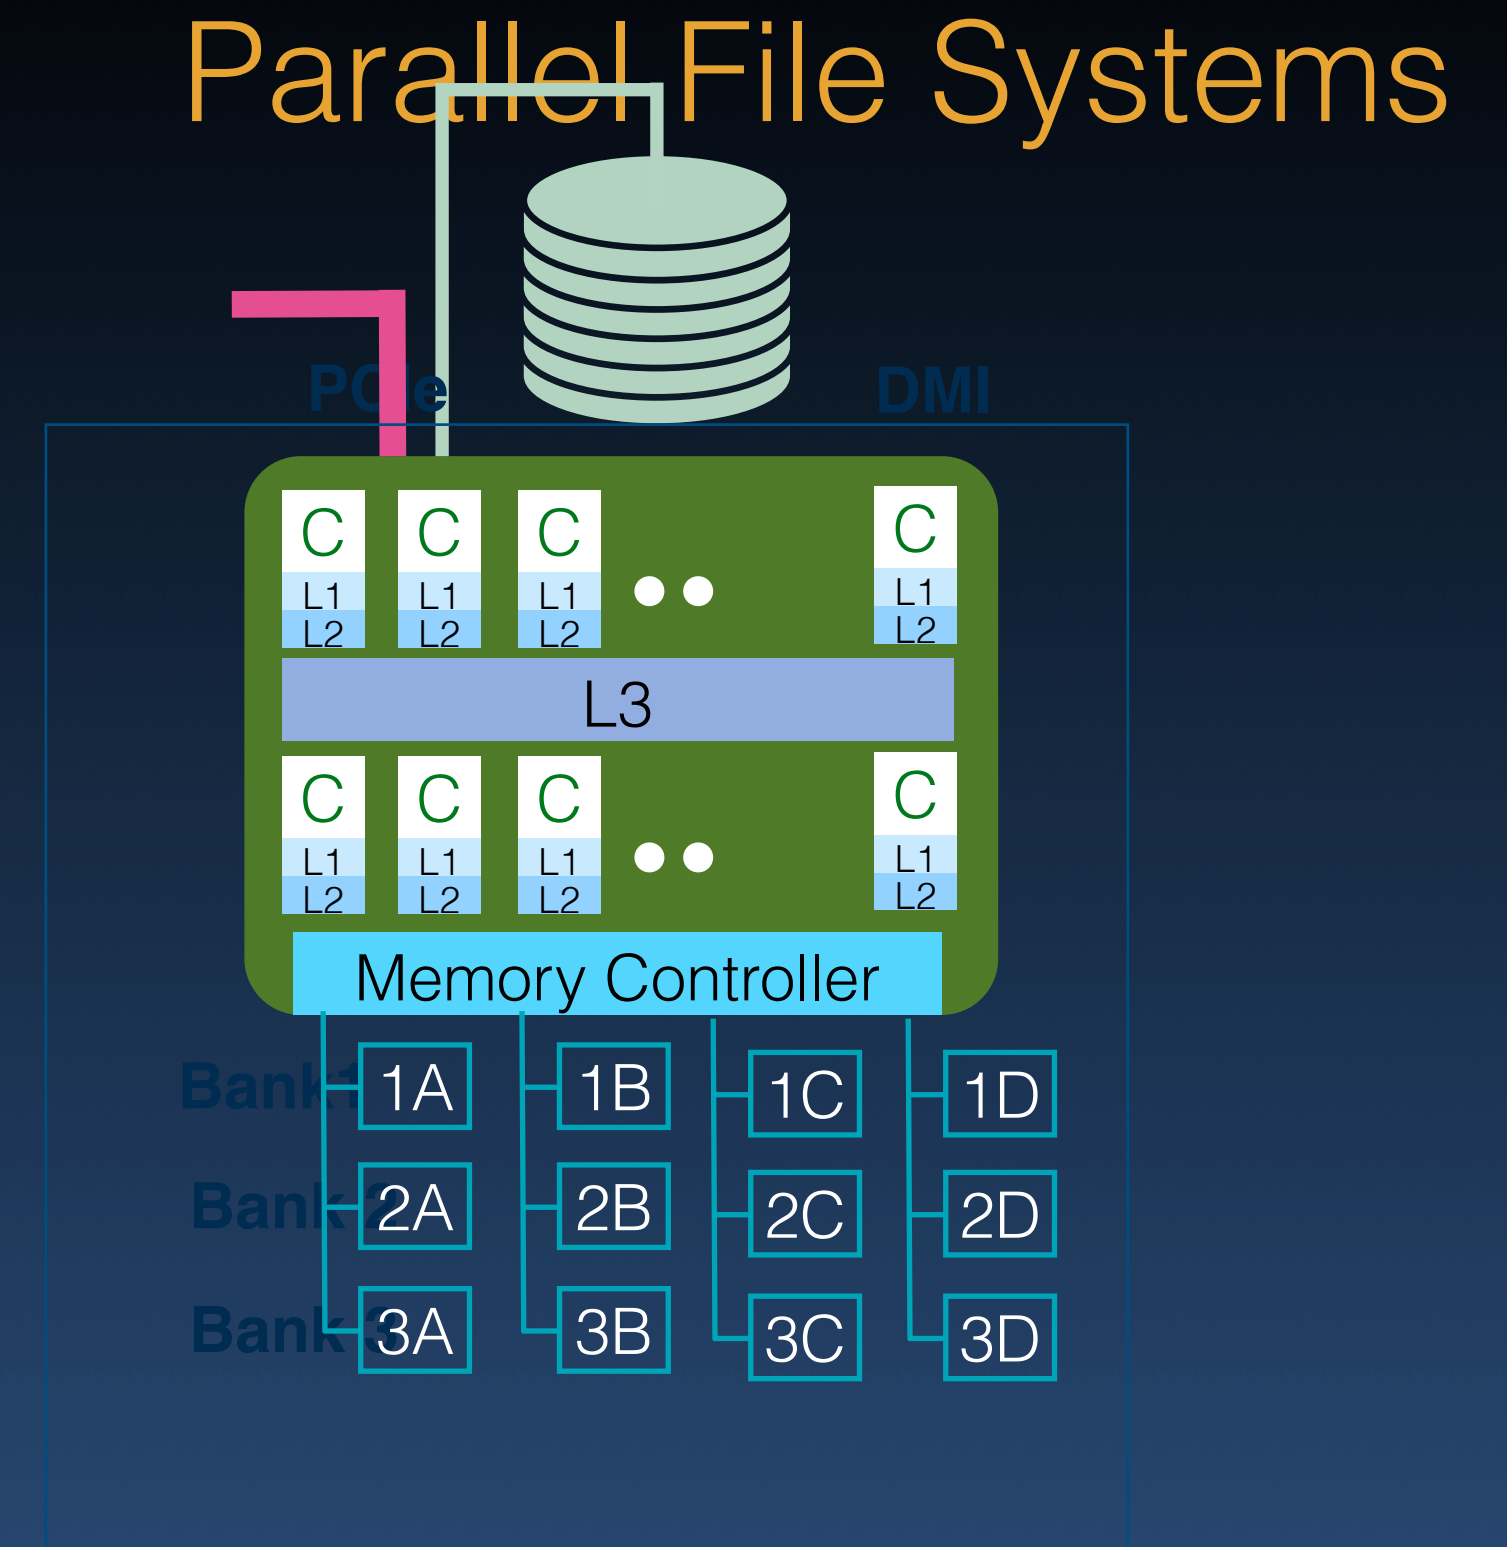


```
MPI_Win_fence(0, windo); // Wait for MPI_Put complete
```

```
if(my_rank == 0)
    use(winbuf+5);
```

- Multiple disk servers
  - ➔ With multiple network paths to disks
- Designed for performance
  - ➔ Large block sizes (~MB)
  - ➔ Parallel fetch
  - ➔ Concurrent I/O
  - ➔ Metadata operations less performant
- Traditional file API
  - ➔ Additional APIs for faster access


- Multiple disk servers
  - ➔ With multiple network paths to disks
- Designed for performance
  - ➔ Large block sizes (~MB)
  - ➔ Parallel fetch
  - ➔ Concurrent I/O
  - ➔ Metadata operations less performant
- Traditional file API
  - ➔ Additional APIs for faster access







# Parallel File Systems

- Multiple disk servers
  - ➔ With multiple network paths to disks
- Designed for performance
  - ➔ Large block sizes (~MB)
  - ➔ Parallel fetch
  - ➔ Concurrent I/O
  - ➔ Metadata operations less performant
- Traditional file API
  - ➔ Additional APIs for faster access



# Parallel File Systems

- Multiple disk servers
  - ➔ With multiple network paths to disks
- Designed for performance
  - ➔ Large block sizes (~MB)
  - ➔ Parallel fetch
  - ➔ Concurrent I/O
  - ➔ Metadata operations less performant
- Traditional file API
  - ➔ Additional APIs for faster access



# PFS Striping

- Configuration per file
  - ➔ number of stripes, stripe size, and OSTs to use


Stripe counts

File A: 3

File B: 1

File C: 1

Stripe size of File C is larger





# PFS Striping

- Configuration per file
  - ➔ number of stripes, stripe size, and OSTs to use

## Stripe counts


File A: 3

File B: 1

File C: 1

```
> lfs getstripe <filename>  
> lfs setstripe <dirname>
```

Stripe size of File C is larger




```
MPI_File fh;  
MPI_Status status;  
  
MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
MPI_Comm_size(MPI_COMM_WORLD, &nprocs);  
bufsize = FILESIZE/nprocs;  
nints = bufsize/sizeof(int);  
MPI_File_open(MPI_COMM_WORLD, "file", MPI_MODE_RDONLY,  
              MPI_INFO_NULL, &fh);  
MPI_File_seek(fh, rank * bufsize, MPI_SEEK_SET);  
MPI_File_read(fh, buf, nints, MPI_INT, &status);  
MPI_File_close(&fh);
```

- 3-tuple: <displacement, etype, filetype>
  - ➔ byte displacement from the start of the file
  - ➔ etype: data unit type
  - ➔ filetype: portion of the file visible to the process
- MPI\_File\_set\_view



```
int MPI_File_set_view(  
    MPI_File fh,  
    MPI_Offset disp,      // in bytes  
    MPI_Datatype etype,    // file's a sequence of etypes  
    MPI_Datatype filetype, // interpret as filetypes  
    char *datarep,  
    MPI_Info info)
```



Derived type




Derived type

etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype
-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------

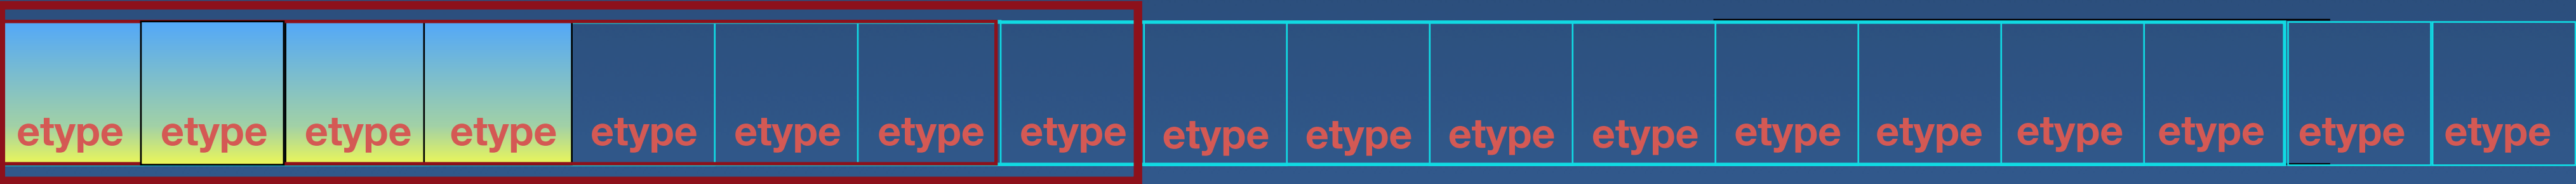

File







etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype	etype
-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------

File



File




File



# File View



etype  
Derived type

filetype



File

# File View



# File

## Example: Write in Set View

```
MPI_File pfile;
```

```
for (i=0; i<BUFSIZE; i++)
```

```
    buf[i] = myrank * BUFSIZE + i;
```

```
MPI_File_open(MPI_COMM_WORLD, "file",  
              MPI_MODE_CREATE | MPI_MODE_WRONLY, MPI_INFO_NULL, &pfile);
```

Blocking, Collective

```
MPI_File_set_view(pfile, myrank * BUFSIZE * sizeof(int), MPI_INT, MPI_INT,  
                  "native", MPI_INFO_NULL);
```

```
MPI_File_write(pfile, buf, BUFSIZE, MPI_INT, MPI_STATUS_IGNORE);
```

```
MPI_File_close(&pfile);
```

Blocking, Individual



- Location

`MPI_File_read_at(fh, offset, buffer, count, datatype, &status)`

- Non-blocking

`MPI_File_iread(fh, buffer, count, datatype, &request)`

- Collective

`MPI_File_read_all(fh, buffer, count, datatype, &status)`

- Shared File pointer (Common data IO)

`MPI_File_read_shared(fh, buffer, count, datatype, &status) // Not collective`

`MPI_File_read_ordered (fh, buffer, count, datatype, &status) // Collective`

- Writes from one process become visible to others at arbitrary times

`MPI_File_set_atomicity (MPI_File fh, int flag);`

→ Collective

`MPI_File_sync ( MPI_File fh );`

→ Collective

→ Flush all writes

## Example: Collective IO

```
MPI_Comm_size(MPI_COMM_WORLD, &size );
MPI_File_open(MPI_COMM_WORLD, "file", MPI_MODE_RDWR|MPI_MODE_CREATE,
              MPI_INFO_NULL, &fh );

MPI_File_write_ordered( fh, buf, 1, MPI_INT, &status );
MPI_Barrier(MPI_COMM_WORLD);           // Let all writes complete


MPI_File_seek( fh, 0, MPI_SEEK_SET );   // Rewind to the top
MPI_File_read_all( fh, buf, size, MPI_INT, &status ); // Everyone reads size ints

MPI_File_seek_shared(fh, 0, MPI_SEEK_SET ); // Rewind to the top again
MPI_File_read_ordered(fh, buf, 1, MPI_INT, &status ); // Read one int in round-robin order

MPI_File_close( &fh );
```



## Example: Views in IO



```
MPI_Comm_rank(MPI_COMM_WORLD, &rank );
MPI_Comm_size(MPI_COMM_WORLD, &size );
MPI_Type_contiguous (4, MPI_DOUBLE, &etype);
MPI_Type_commit ( &etype );
```

```
for ( i = 0; i < 4; i++) {
    displ[i] = rank + i * size;
    blocklength[i] = 1;
}
MPI_Type_indexed (4, blocklength, displ, etype, &filetype );
MPI_Type_commit ( &filetype );
```

```
MPI_File_open ( MPI_COMM_WORLD,"file", MPI_MODE_RDONLY, MPI_INFO_NULL , &fh);
MPI_File_set_view (fh, 0, etype, filetype, "native", MPI_INFO_NULL);
MPI_File_read_all (fh, buf, 16, etype, &status );
MPI_File_close( &fh );
```


## Example: Views in IO

```
MPI_Comm_rank(MPI_COMM_WORLD, &rank );  
MPI_Comm_size(MPI_COMM_WORLD, &size );  
MPI_Type_contiguous (4, MPI_DOUBLE, &etype);  
MPI_Type_commit ( &etype );
```

```
for ( i = 0; i < 4; i++) {  
    displ[i] = rank + i * size;  
    blocklength[i] = 1;  
}
```

```
MPI_Type_indexed (4, blocklength, displ, etype, &filetype );  
MPI_Type_commit ( &filetype );
```

```
MPI_File_open ( MPI_COMM_WORLD,"file", MPI_MODE_RDONLY, MPI_INFO_NULL , &fh);  
MPI_File_set_view (fh, 0, etype, filetype, "native", MPI_INFO_NULL);  
MPI_File_read_all (fh, buf, 16, etype, &status );  
MPI_File_close( &fh );
```




## Example: Views in IO

```
MPI_Comm_rank(MPI_COMM_WORLD, &rank );
MPI_Comm_size(MPI_COMM_WORLD, &size );
MPI_Type_contiguous (4, MPI_DOUBLE, &etype);
MPI_Type_commit ( &etype );
```

```
for ( i = 0; i < 4; i++) {
    displ[i] = rank + i * size;
    blocklength[i] = 1;
}
```

```
MPI_Type_indexed (4, blocklength, displ, etype, &filetype );
MPI_Type_commit ( &filetype );
```

```
MPI_File_open ( MPI_COMM_WORLD,"file", MPI_MODE_RDONLY, MPI_INFO_NULL , &fh);
MPI_File_set_view (fh, 0, etype, filetype, "native", MPI_INFO_NULL);
MPI_File_read_all (fh, buf, 16, etype, &status );
MPI_File_close( &fh );
```






## Example: Views in IO

```
MPI_Comm_rank(MPI_COMM_WORLD, &rank );  
MPI_Comm_size(MPI_COMM_WORLD, &size );  
MPI_Type_contiguous (4, MPI_DOUBLE, &etype);  
MPI_Type_commit ( &etype );
```

```
for ( i = 0; i < 4; i++) {  
    displ[i] = rank + i * size;  
    blocklength[i] = 1;  
}
```

```
MPI_Type_indexed (4, blocklength, displ, etype, &filetype );  
MPI_Type_commit ( &filetype );
```

```
MPI_File_open ( MPI_COMM_WORLD,"file", MPI_MODE_RDONLY, MPI_INFO_NULL , &fh);  
MPI_File_set_view (fh, 0, etype, filetype, "native", MPI_INFO_NULL);  
MPI_File_read_all (fh, buf, 16, etype, &status );  
MPI_File_close( &fh );
```



## Example: Views in IO

```
MPI_Comm_rank(MPI_COMM_WORLD, &rank );  
MPI_Comm_size(MPI_COMM_WORLD, &size );  
MPI_Type_contiguous (4, MPI_DOUBLE, &etype);  
MPI_Type_commit ( &etype );
```

```
for ( i = 0; i < 4; i++) {  
    displ[i] = rank + i * size;  
    blocklength[i] = 1;  
}
```

```
MPI_Type_indexed (4, blocklength, displ, etype, &filetype );  
MPI_Type_commit ( &filetype );
```

```
MPI_File_open ( MPI_COMM_WORLD,"file", MPI_MODE_RDONLY, MPI_INFO_NULL , &fh);  
MPI_File_set_view (fh, 0, etype, filetype, "native", MPI_INFO_NULL);  
MPI_File_read_all (fh, buf, 16, etype, &status );  
MPI_File_close( &fh );
```

