

COL380

Introduction to
Parallel & Distributed Programming

- EREW (Exclusive Read Exclusive Write)
 - ➔ Only one processors may read or write any given location in a step
- CREW (Concurrent Read Exclusive Write)
 - ➔ Many processors can simultaneously read a location, but only one may write
- CRCW (Concurrent Read Concurrent Write)
 - ➔ Many processors can read/write the same memory location
- ERCW (Exclusive Read Concurrent Write)
 - ➔ Not commonly used

Concurrent Write (CW)

- **Priority CW**
 - ➔ Higher priority processor (normally lower index) wins
- **Common CW**
 - ➔ Succeeds only if all writes have the same value
- **Arbitrary/Random CW**
 - ➔ One of the values is randomly chosen

Concurrent Write (CW)

- **Priority CW**
 - ➔ Higher priority processor (normally lower index) wins
- **Common CW**
 - ➔ Succeeds only if all writes have the same value
- **Arbitrary/Random CW**
 - ➔ One of the values is randomly chosen

EREW ≤ CREW ≤ Common ≤ Arbitrary ≤ Priority

Less powerful

More powerful

Parallel Addition

$p = n; B[i] = A[i]$

p0

p1

p2

p3

p4

p5

p6

p7

$p = p/2; \text{if}(i < p) B[i] = B[2i] + B[2i+1]$

p0

p1

p2

p3

$p = p/2; \text{if}(i < p) B[i] = B[2i] + B[2i+1]$

p0

p1

$p = p/2; \text{if}(i < p) B[i] = B[2i] + B[2i+1]$

p0

```
p = n/2  
forall i < n  
    B[i] = A[i]  
while(p > 0) {  
    forall i < p  
        B[i] = B[2i] + B[2i+1]  
    p = p/2;  
}
```

(assumes *n* is a power of 2)

- processors: *n*
- time: $O(\log n)$
- Speed-up: $n/(\log n)$
- Efficiency: $1/\log(n)$
- Cost: $n \log n$
- Work: *n*

Linear Search

$p < n$

- n input integers in n memory cells
- Does x exist in the input?
 - x is initially stored in shared memory

Algorithm

step1: If p_0 , broadcast x

step2: Processor p_i : search in i th $[n/p]$ block and {set flag f_i }

step3: If p_0 , check if 'any' flag is 1 and print answer

EREW

• $\log(p)$

• n/p

• $\log(p)$

CREW

• 1

• n/p

• $\log(p)$

CRCW

• 1

• n/p

• 1

Small Processor Count

- **Lemma**

- ➔ Any problem that can be solved on a p -processor PRAM in t steps can be solved on a p' -processor PRAM in $t' = O(t * p / p')$ steps (assuming the same size of shared memory).

- **Proof**

- ➔ Partition p *simulated* processors into p' groups of size p/p' each.

- ➔ Allocate each group to one *available* processor

- ➔ An available processor executes on behalf of its group, one step at a time

- Execute all READ substeps

- Execute all LOCAL computation substeps

- Execute all WRITE substeps

- ➔ Assumes local register space

- **Lemma**

- ➡ Any problem that can be solved on a p -processor and m -cell PRAM in t steps can be solved on a $\max(p, m')$ -processor and m' -cell PRAM in $O(tm/m')$ steps.

- **Proof:**

- ➡ Partition m simulated cells into m' segments S_i , each of size m/m'
 - ➡ Each available processor P'_i simulates processor P_i , $i = [1..p]$
 - ➡ Each available processor P'_i , $i = [1..m']$, stores S_i in its local memory and uses $M'[i]$ as a public cell to simulate accesses to cells of S_i
 - ➡ To simulate **READ**, P'_i , $i = [1..\max(p, m')]$, repeats, for $k = [1..m/m']$:
 - write the value of the k -th cell of S_i into $M'[i]$, $i = [1..m']$
 - read the value if the original processor P_i , $i = [1..p]$, would read that cell
 - ➡ P'_i then simulates the computation substep of P_i , $i = [1..p]$, in one step

- **Lemma**

- ➡ Any problem that can be solved on a p -processor and m -cell CREW PRAM in t steps can be solved on a p -processor and m -cell EREW PRAM in $O(tp)$ steps.

- **Proof:**

- ➡ Complete each READ substep in p rounds
 - ➡ In round i , processor P_i executes its READ

Possible in fewer steps?

Performance Evaluation

- Two parameters

- $p(n)$, $t(n)$

- Generally, use work, $W(n)$

- If $W(n)$ similar, use $t(n)$

- Speedup/Scalability

- Absolute: over best sequential algorithm

- Relative: over the 1-processor implementation of the same algorithm

- Work-optimal \Rightarrow work = $O(\text{serial complexity})$

- $p(n)$ is hidden but important

- ▶ $W_1(n) = O(n)$; $t_1(n) = O(n)$

- ▶ $W_2(n) = O(n \log n)$ and $t_2(n) = O(\log n)$

Work Time Scheduling Principle

- Design algorithm in terms of
 - Total work done per 'time step': $W_i(n)$
 - $t(n)$ steps
- Total work done $W(n) = \sum W_i(n)$
- For each time step i :
 - divide the work $W_i(n)$ among p processors
 - ▶ Time $\leq \sum \lceil W_i(n)/p \rceil \leq \lceil W(n)/p \rceil + t(n)$
- Cost = $t(n,p) * p$

Work \leq Cost. Cost optimality is more stringent.

Brent's Theorem

- Time taken by p processors:
 - $t(n, p) = O(W(n)/p + t(n))$
- Cost = $p * t(n, p) = O(W(n) + p * t(n))$
- Work = Cost if:
 - $W(n) + p * t(n) = O(W(n))$
 - Or, $p = O(W(n)/t(n))$

Optimal Summation

- Using n processors for parallel sum
 - $\text{Work} = O(n)$, $\text{Cost} = O(n \log n)$
- Suppose, we use $n/\log n$ processors
 - forall: Compute sum of its share of $\log n$ local elements
 - Compute sum of $n/\log n$ partial sums
 - ▶ Using previous parallel sum algorithm (with $n/\log n$ processors)
- $\text{Time} = \log n + \log(n/\log n) = \log(n)$
- $\text{Cost} = n/\log n * \log n = n$

Favor long "sequential" sections

- If **sequentially** optimal algorithm is $O(t'(n))$

→ Work done by **Work-optimal** parallel algorithm:

▶ $O(t'(n))$ (with time $t(n)$).

→ Work-scheduling on p processors takes time:

▶ $t(n, p) = O(t'(n)/p + t(n))$

→ Optimal speed-up: $t'(n)/t(n, p) = \theta(p)$, if

▶ $[p \cdot t'(n)] / [t'(n) + p \cdot t(n)] = \theta(p)$

- **Work-time** optimal if:

→ $t(n)$ cannot be improved