

COL380

Introduction to  
Parallel & Distributed Programming

## Why PRAM?

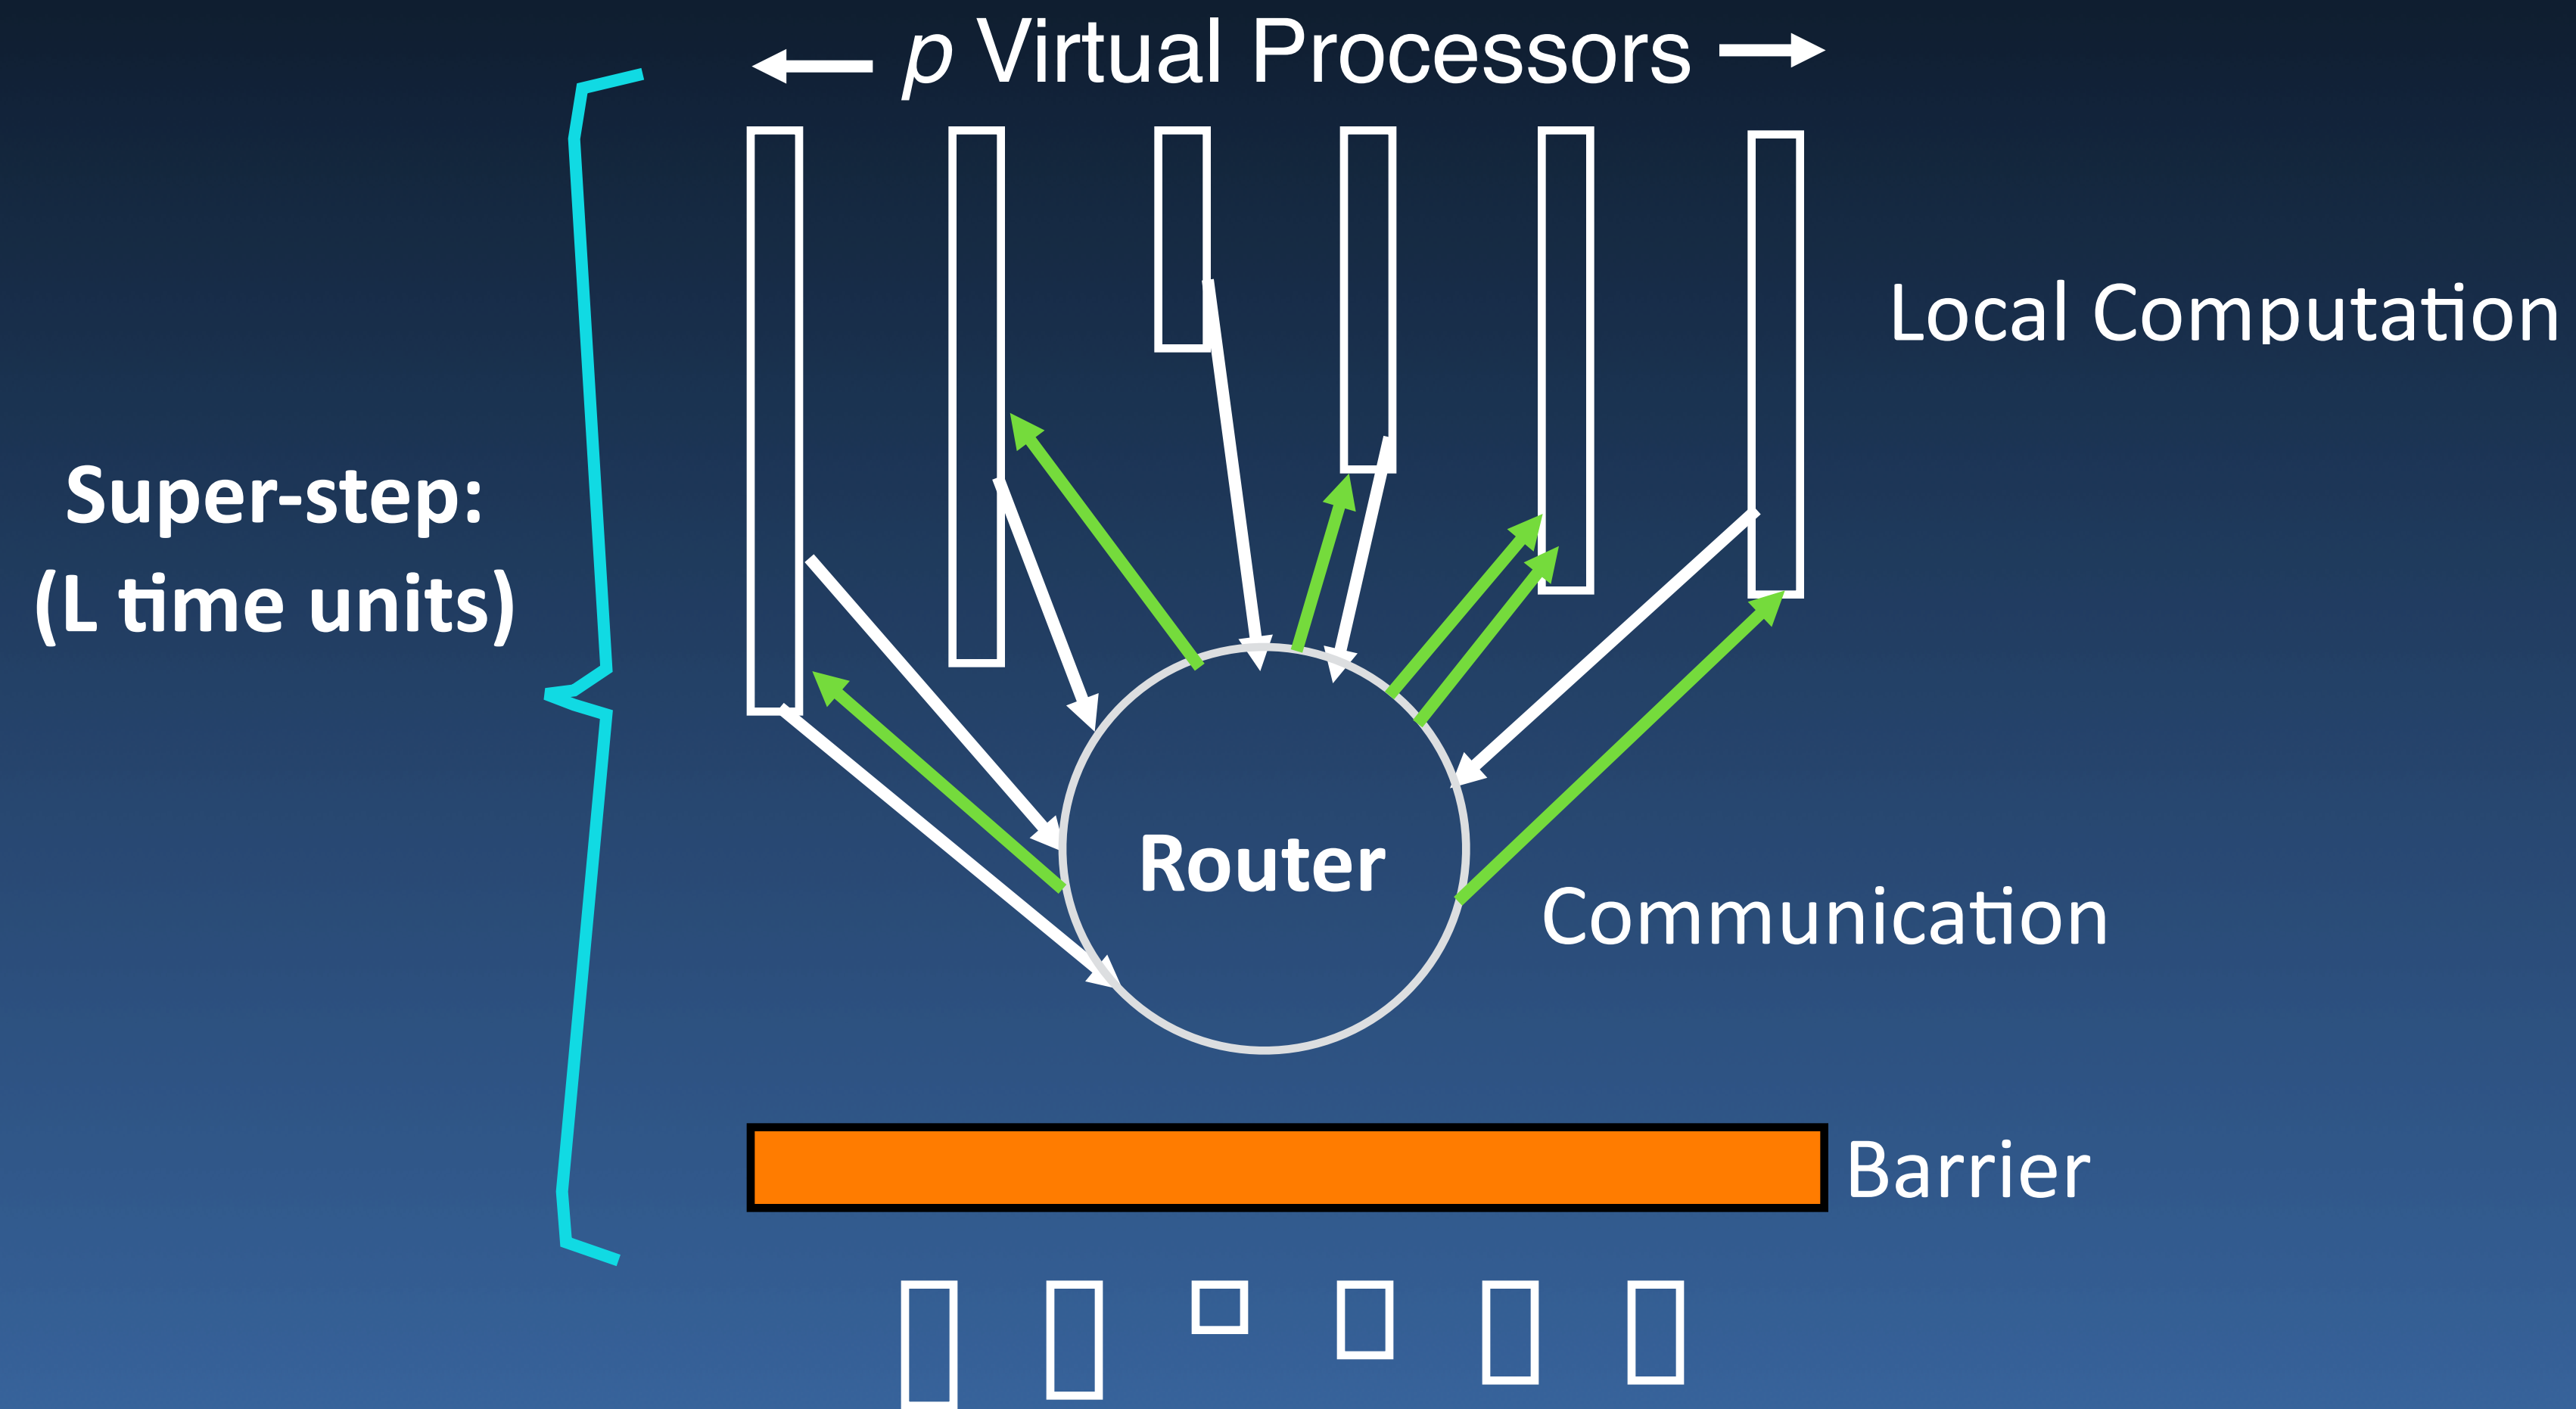
- Easy to design, specify, analyze algorithms
  - ➔ Independent of machine details
- Fidelity of predicted performance
  - ➔ Not many surprises for shared-memory architecture
  - ➔ Only partially successful for distributed memory
  - ➔ Note that memory-access and message latency is often bounded
- Strong model
  - ➔ Possible to simulate on a wide variety of hardware
  - ➔ Poor PRAM solution often implies a hard problem

- Fine-grained synchronization

## Bulk Synchronous Parallel Model

- A set of virtual (processor, memory) pairs
  - ➔ No notion of “locality” in mapping to physical processor
- A point-to-point interconnect
- Barrier synchronization (all or subset)
- Repeat “super-step”:
  - ➔ Local computation
  - ➔ Communication
  - ➔ Barrier synchronization

Designed as a Bridging  
model for parallel  
programming





- Steps are machine wide, similar to PRAM
  - ➔ Communication and computation
- Simple to program and analyze algorithms
  - ➔ Do not have to separately model individual communication
- “Bridge” between PRAM and real architecture
- Locality of processor or computation ignored
  - ➔ Leaves such optimization out of the equation

- In each super-step, VP does a task with
  - ➔ implicit message arrivals, local computation, message sends
  - ➔ Computation may only access data available at the beginning of super-step
- VPs synchronize at regular intervals of  $L$  time units
  - ➔  $L$ : the periodicity parameter
    - ▶  $L$  may be a constant
    - ▶ If a VP not done with in  $L$ , the super-step continues (a multiple of  $L$ )

- Realize arbitrary  $h$ -relation in a super-step:
  - VPs send (receive) at most  $h$  messages
    - ▶ each of length “1”
  - Communication cost:  $gh = th + s$ 
    - ▶  $t$ : network latency
    - ▶  $s$ : constant overhead
  - Ignores actual transfer size, distribution

- Barrier cost of  $l$ 
  - May depend on the count of virtual processors,  $p$
  - A lower bound on  $l$  is the diameter of the network
- Super-step cost
  - cost of the longest running local computation
  - cost of communication
  - cost of the barrier synchronization
  - $w + hg + l$  (Non Overlap model)
- Total cost = summation over super-steps  $s$ :
  - $\sum w_s + g \sum h_s + Nl$  ( $N$  is the number of super-steps)



## Complexity Class NC

- $t(n, p(n)) = O(\log_{k_1} n)$  with  $p(n) = O(n^{k_2})$  processors

→  $k_1, k_2 = O(1)$

→ Assume PRAM model (Doesn't matter which)

→ Robust with respect to other models as well

Is parallel “p-ary” search in NC?

- Many NC algorithms are work optimal

→ e.g., for a linear-time sequential algorithm:

▶  $t(n, p(n)) = O(\log n)$  with  $p(n) = n/\log n$

▶  $t(n, p(n)) = O(\log \log n)$  with  $p(n) = n/\log \log n$

$NC \subseteq P$   
Open Question: is  $P == NC$ ?

- “Inherently sequential,” if no NC algorithms exist for a problem