


COL380

Introduction to
Parallel & Distributed Programming

**Parallel algorithm technique:
Balanced binary tree**


Reduction

- n operands $\Rightarrow \log n$ steps
- How do you map?
 - $n/2^i$ processors per step
 - step i : if $!(id \% 2^i)$




Reduction

- n operands $\Rightarrow \log n$ steps
- Only have p processors per step
- Agglomerate and Map



Prefix Sums


- $P[0] = x[0]$
- For $i = 1$ to $n-1$
 - $P[i] = P[i-1] + x[i]$



Recursive Prefix Sums

- $P[0] = x[0]$
- For $i = 1$ to $n-1$
 - $P[i] = P[i-1] + x[i]$

$$T(n) = T(n/2) + O(1)$$
$$W(n) = 2W(n/2) + Kn/2$$



Recursive Prefix Sums

- $P[0] = x[0]$
- For $i = 1$ to $n-1$
 - $P[i] = P[i-1] + x[i]$

$$T(n) = T(n/2) + O(1)$$
$$W(n) = 2W(n/2) + Kn/2$$

$$W(n) = O(n \log n)$$




Recursive Prefix Sums

- $P[0] = x[0]$
- For $i = 1$ to $n-1$
 - $P[i] = P[i-1] + x[i]$

$$T(n) = T(n/2) + O(1)$$
$$W(n) = 2W(n/2) + Kn/2$$

$$W(n) = O(n \log n)$$




Recursive Prefix Sums

$$T(n) = T(n/2) + O(1)$$
$$W(n) = 2W(n/2) + Kn/2$$

$$W(n) = O(n \log n)$$

- $P[0] = x[0]$
- For $i = 1$ to $n-1$
 - $P[i] = P[i-1] + x[i]$



Recursive Prefix Sums


- $P[0] = x[0]$
- For $i = 1$ to $n-1$
 - $P[i] = P[i-1] + x[i]$

$$T(n) = T(n/2) + O(1)$$
$$W(n) = 2W(n/2) + Kn/2$$

$$W(n) = O(n \log n)$$

$$T(n) = T(n/2) + O(1)$$
$$W(n) = W(n/2) + Kn/2$$

$$W(n) = O(n)$$




Recursive Prefix Sums

```
prefixSums(P, x, [0:n])  
{  
  forall i in [0:n/2)  
    y[i] = OP(x[2*i], x[2*i+1])  
  prefixSum(z, y, [0:n/2))  
  P[0] = x[0]  
  forall i in [1:n)  
    if(i&1) P[i] = z[i/2]  
    else   P[i] = OP(z[i/2-1 ], x[i])  
}
```

Or $OP^{-1}(z[i/2], x[i])$,
if op invertible


Prefix Sum



Prefix Sums (Binary Tree)

$$P[0] = x[0]$$

For $i = 1$ to $n-1$


$$P[i] = P[i-1] + x[i]$$


Prefix Sums (Binary Tree)

$$P[0] = x[0]$$


For $i = 1$ to $n-1$

$$P[i] = P[i-1] + x[i]$$



Prefix Sums (Binary Tree)


$P[0] = x[0]$
For $i = 1$ to $n-1$
 $P[i] = P[i-1] + x[i]$



Prefix Sums (Binary Tree)

$P[0] = x[0]$
 For $i = 1$ to $n-1$
 $P[i] = P[i-1] + x[i]$

forall $i = 0$ to n
 $B[0][i] = A[i]$
 for $h = 1$ to $\log n$
forall i in $0:n/2^h$
 $B[h][i] = B[h-1][2i] \text{ OP } B[h-1][2i+1]$
 for $h = \log n$ to 0
 $C[h][0] = B[h][0]$
forall i in $1:n/2^h$
 Odd i : $C[h][i] = C[h+1][i/2]$
 Even i : $C[h][i] = C[h+1][i/2-1] \text{ OP } B[h][i]$



Balanced Tree Approach

- Build binary tree on the input
- Hierarchically divide into groups
 - and groups of groups..
- Traverse tree upwards/downwards
- Useful to think of “tree” network topology
 - Only for algorithm design
 - Later map sub-trees to processors

Parallel algorithm techniques:


PARTITIONING

Merge Sorted Sequences (A,B)

- Determine Rank of each element in $A \cup B$
- $\text{Rank}(x, A \cup B) = \text{Rank}(x, A) + \text{Rank}(x, B)$
 - Only need to compute the rank in the other list, if A and B are each sorted already
- Find $\text{Rank}(A, B)$, and similarly $\text{Rank}(B, A)$
- Find Rank by binary search
 - $O(\log n)$ time
- $O(n \log n)$ work


Towards Optimal Merge (A,B)

- Partition A and B into $\log n$ sized blocks
- Select from B, elements $i * \log n$, $i \in 0:n/\log n$
- Rank each selected element of B in A
 - Binary search
- Merge pairs of sub-sequences
 - If $|A_i| = \log(n)$, Sequential merge in time $O(\log(n))$
 - Otherwise, partition A_i into $\log n$ blocks
 - And Recursively subdivide B_i into sub-sub-sequences




Towards Optimal Merge (A,B)

- Partition A and B into $\log n$ sized blocks
- Select from B, elements $i * \log n$, $i \in 0:n/\log n$
- Rank each selected element of B in A
 - Binary search
- Merge pairs of sub-sequences
 - If $|A_i| = \log(n)$, Sequential merge in time $O(\log(n))$
 - Otherwise, partition A_i into $\log n$ blocks
 - And Recursively subdivide B_i into sub-sub-sequences



Towards Optimal Merge (A,B)

- Partition A and B into $\log n$ sized blocks
 - Select from B, elements $i * \log n$, $i \in 0:n/\log n$
 - Rank each selected element of B in A
 - Binary search
 - Merge pairs of sub-sequences
 - If $|A_i| = \log(n)$, Sequential merge in time $O(\log(n))$
 - Otherwise, partition A_i into $\log n$ blocks
 - And Recursively subdivide B_i into sub-sub-sequences
- Total time is $O(\log(n))$
 - Total work is $O(n)$



Can we do better?


Fast Merge (A,B)

- Partition A and B into \sqrt{n} blocks
- Select from B, elements $i\sqrt{n}$, $i \in (0: \sqrt{n}]$
- Rank each selected element of B in A
 - Parallel search using \sqrt{n} processors each search
- Similarly rank \sqrt{n} selected elements from A in B
- Recursively merge pairs of sub-sequences
 - Total time: $T(n) = O(1) + T(\sqrt{n}) = O(\log \log n)$
 - Total work: $W(n) = O(n) + \sqrt{n} W(\sqrt{n}) = O(n \log \log n)$
- “Fast” but still need to reduce *work*

Not work optimal


Optimal Merge (A,B)

- Use the fast, but non-optimal, algorithm on small enough subsets
- Subdivide A and B into blocks of size $\log \log n$
 - A_1, A_2, \dots
 - B_1, B_2, \dots
- Select first element of each block
 - $A' = p_1, p_2, \dots$
 - $B' = q_1, q_2, \dots$
- Now merge $\log \log n$ sized blocks $n/\log \log n$ times




Optimal Merge (A,B)

- Use the fast, but non-optimal, algorithm on small enough subsets
- Subdivide A and B into blocks of size $\log \log n$
 - A_1, A_2, \dots
 - B_1, B_2, \dots
- Select first element of each block
 - $A' = p_1, p_2, \dots$
 - $B' = q_1, q_2, \dots$
- Now merge $\log \log n$ sized blocks $n/\log \log n$ times




Optimal Merge (A,B)

- Use the fast, but non-optimal, algorithm on small enough subsets
- Subdivide A and B into blocks of size $\log \log n$
 - A_1, A_2, \dots
 - B_1, B_2, \dots
- Select first element of each block
 - $A' = p_1, p_2, \dots$
 - $B' = q_1, q_2, \dots$
- Now merge $\log \log n$ sized blocks $n/\log \log n$ times




Optimal Merge (A,B)

- Use the fast, but non-optimal, algorithm on small enough subsets
- Subdivide A and B into blocks of size $\log \log n$
 - A_1, A_2, \dots
 - B_1, B_2, \dots
- Select first element of each block
 - $A' = p_1, p_2, \dots$
 - $B' = q_1, q_2, \dots$
- Now merge $\log \log n$ sized blocks $n/\log \log n$ times




Optimal Merge (A,B)

- Use the fast, but non-optimal, algorithm on small enough subsets
- Subdivide A and B into blocks of size $\log \log n$
 - A_1, A_2, \dots
 - B_1, B_2, \dots
- Select first element of each block
 - $A' = p_1, p_2, \dots$
 - $B' = q_1, q_2, \dots$
- Now merge $\log \log n$ sized blocks $n/\log \log n$ times



Optimal Merge (A,B)

1. Merge A' and B' – find $\text{Rank}(A':B')$, $\text{Rank}(B':A')$
 - using fast non-optimal algorithm
 - Time = $O(\log \log n)$
 - Work = $O(n)$
2. Compute $\text{Rank}(A':B)$ and $\text{Rank}(B':A)$
 - If $\text{Rank}(p_i, B)$ is r_i , p_i lies in block B_{r_i}
 - Sequentially
 - Time = $O(\log \log n)$
 - Work = $O(n)$
3. Compute ranks of remaining elements
 - Sequentially
 - Time = $O(\log \log n)$
 - Work = $O(n)$



Quick Sort

- Choose the pivot
 - Select median?
- Subdivide into two groups
 - Group sizes linearly related with high probability (expect $\log(n)$ rounds)
- Sort each group independently
- Time per round = $O(\log n)$
- Work per round = $O(n)$

```
QuickSort(int A[], int first, int last)
{
    Select random m in [first:last] // A[m] is pivot
    forall i in [first:last]
        flag[i] = A[i] < A[m];
    Split(A); // Separate flag values 0 and 1
               // Prefix sum, A[P[m]] = A[m]
    Quicksort A[first:k-1] and A[k+1:last]
}
```

$$T(n) \sim T(n/2) + O(\log n)$$
$$W(n) \sim 2W(n/2) + O(n)$$

Partitioning Approach

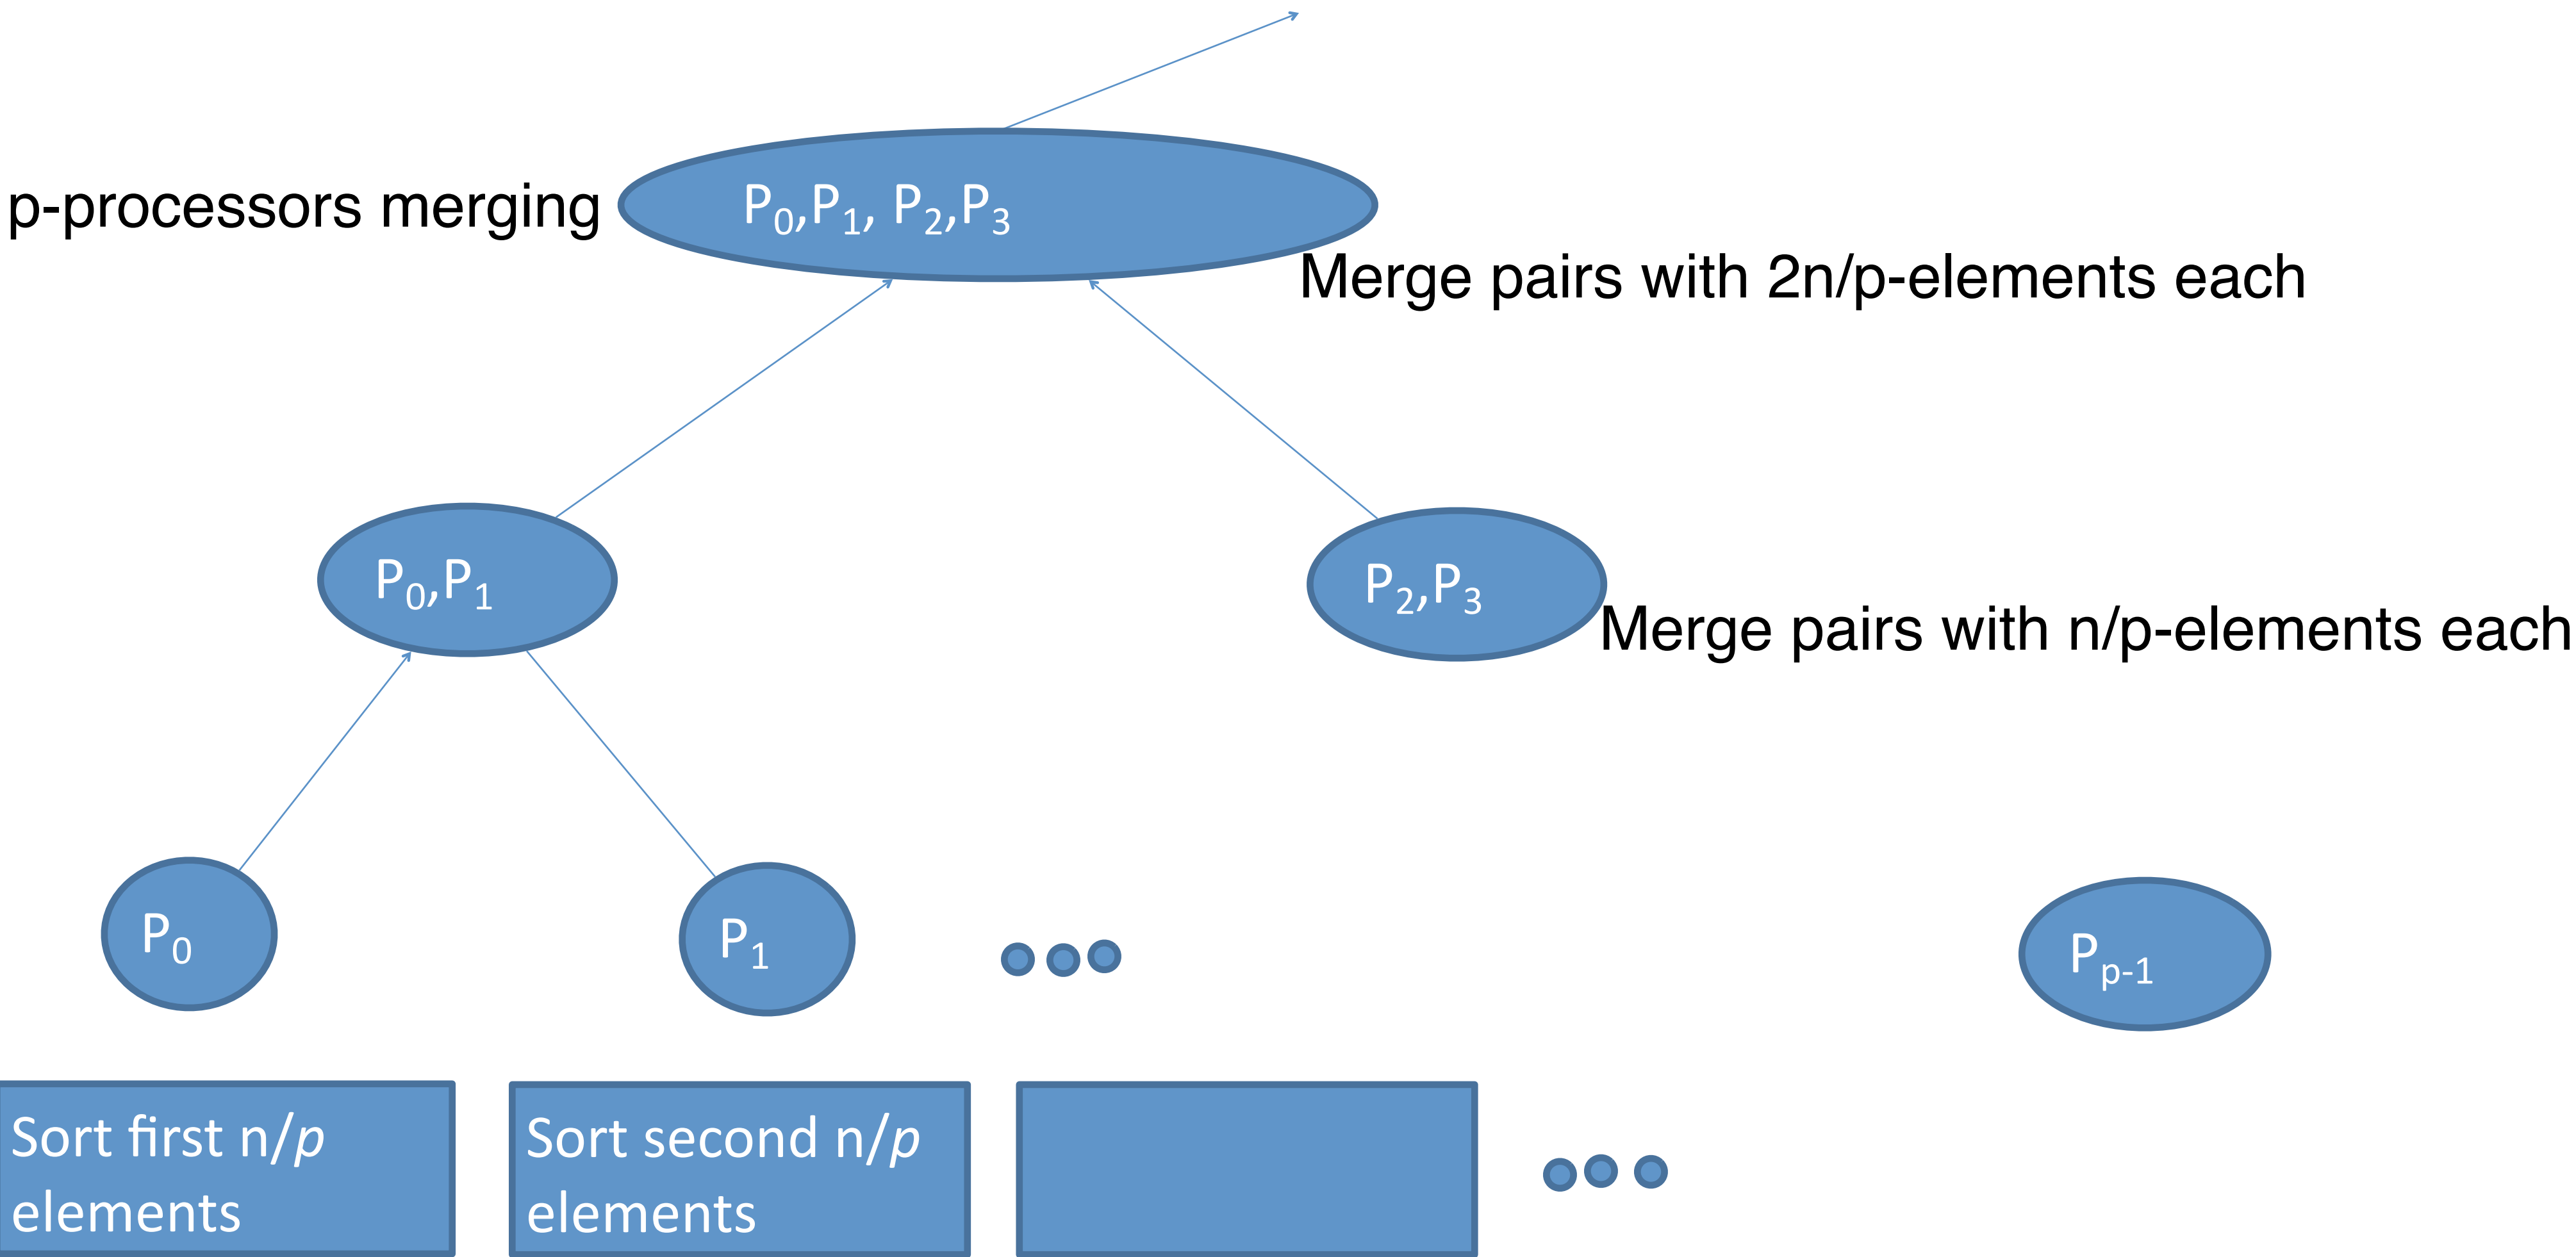
- Break into p roughly equal sized problems
- Solve each sub-problem
 - Preferably, independently of each other
- Focus on subdividing into independent parts

Parallel algorithm techniques: DIVIDE AND CONQUER

Merge Sort

- Partition data into two halves
 - Assign half the processors to each half
 - If only one processor remains, sequentially sort
- Sort each half
- Merge results
- $T(n) = T(n/2) + O(\log \log n)$
- $W(n) = 2 W(n/2) + O(n)$

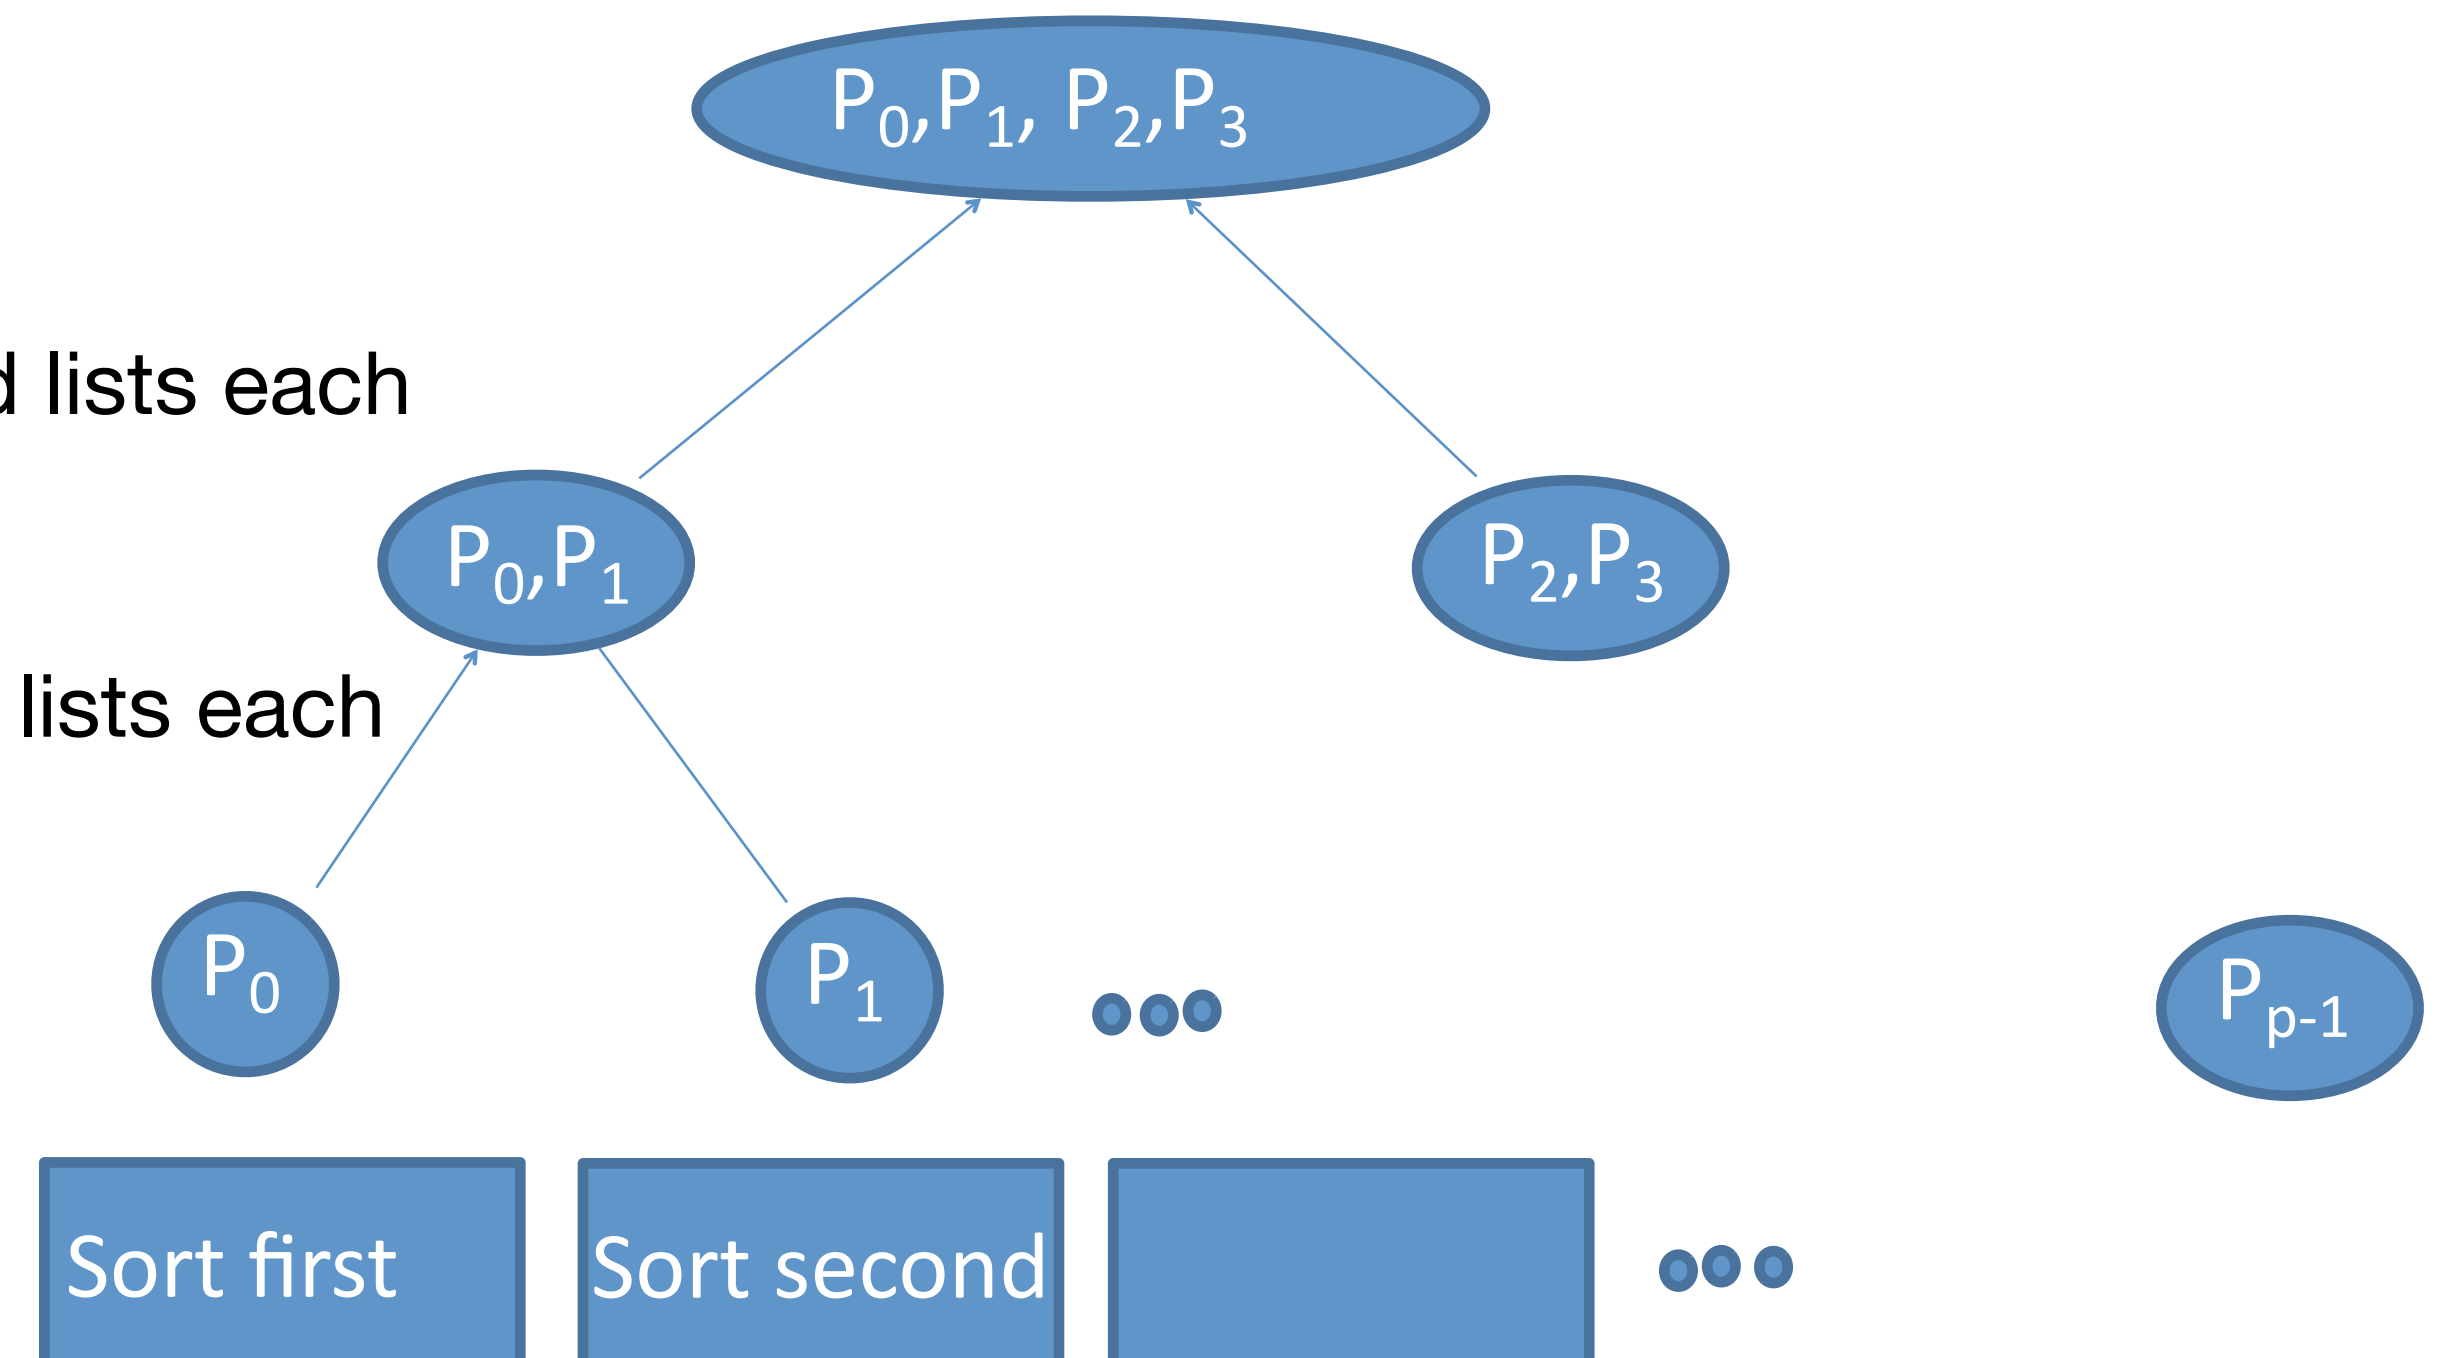
Sort n/p elements, then Merge



HOW EFFICIENTLY CAN YOU MERGE?


Merge Sort

- Divide into p groups
 - Locally sort each group
 - $n/p \log (n/p) = O((n \log n)/p)$
- Parallel merge p groups
 - Binary tree: $\log(p)$ stages
 - @leaf (level 1), 2 processors merge two n/p sized lists each
 - time = $O(n/p)$
 - @Level i : 2^{i+1} processors merge two $2^i n/p$ sized lists each
 - time = $O(n/p)$
 - @Root: p processor merge two $n/2$ sized list
 - time = $O(n/p)$
 - $O(n/p) \log p$



Hyper Quick Sort

- Partition into p groups
 - Sort each group independently
 - $O(n/p \log n)$
- Choose median of one of the groups
- Partition each group into “less” and “more” set
 - Binary search of the median: $(\log n)$
- Separate into low and high
 - Merge $p/2$ “less” and $p/2$ “more” pairs
 - Each sequentially: $O(n/p)$
- Now we have $p/2$ “less” lists and $p/2$ “more” lists
 - Partition and recurse
- Total time = $O(n/p \log n)$ with high probability



Parallel Bucket Sort

- Divide the range $[a,b]$ of numbers into p equal sub-ranges
 - Or, buckets
- Divide input into p blocks
 - arbitrarily
- Each p_j sorts the elements in its block into p buckets
 - “Sends” i^{th} bucket to p_i
 - p_i collects bucket i from each other processor
 - For uniformly distributed input, expected bucket size is uniform
- Locally sort each bucket

$O(n/p \log n/p + p \log p)$?



But, real risk of load imbalance

Sample sort:

Choose a sample of size s

Sort the samples

Choose **B-1** evenly spaced element from the sorted list

These splitters provide ranges for **B** buckets

Parallel Splitter Selection

- Divide n elements equally into B blocks
- (Quick)Sort each block
- For each sorted block:
 - Choose $B-1$ evenly spaced samples
- Use the $B*(B-1)$ elements as samples
- Sort the samples
 - Choose $B-1$ **Splitters**
- Arrange elements by bucket in output array
 - Count the number of elements in each bucket
 - Perform prefix Sum of counts; Reserve space per bucket
 - In-place
 - No bucket contains more than $2*n/B$ elements

$(n/B \log n/B)$

$(B^2 \log B)$

$(n/B + B \log B)$

Parallel algorithm techniques: ACCELERATED CASCADING

Min-find

Input: array with n numbers

Algorithm A1 using $O(n^2)$ processors:

parallel for i in $(0:n]$

$M[i] := 0$

parallel for i, j in $(0:n]$

if $i \neq j \ \&\& \ C[i] < C[j]$

$M[j] = 1$

parallel for i in $(0:n]$

if $M[i] = 0$

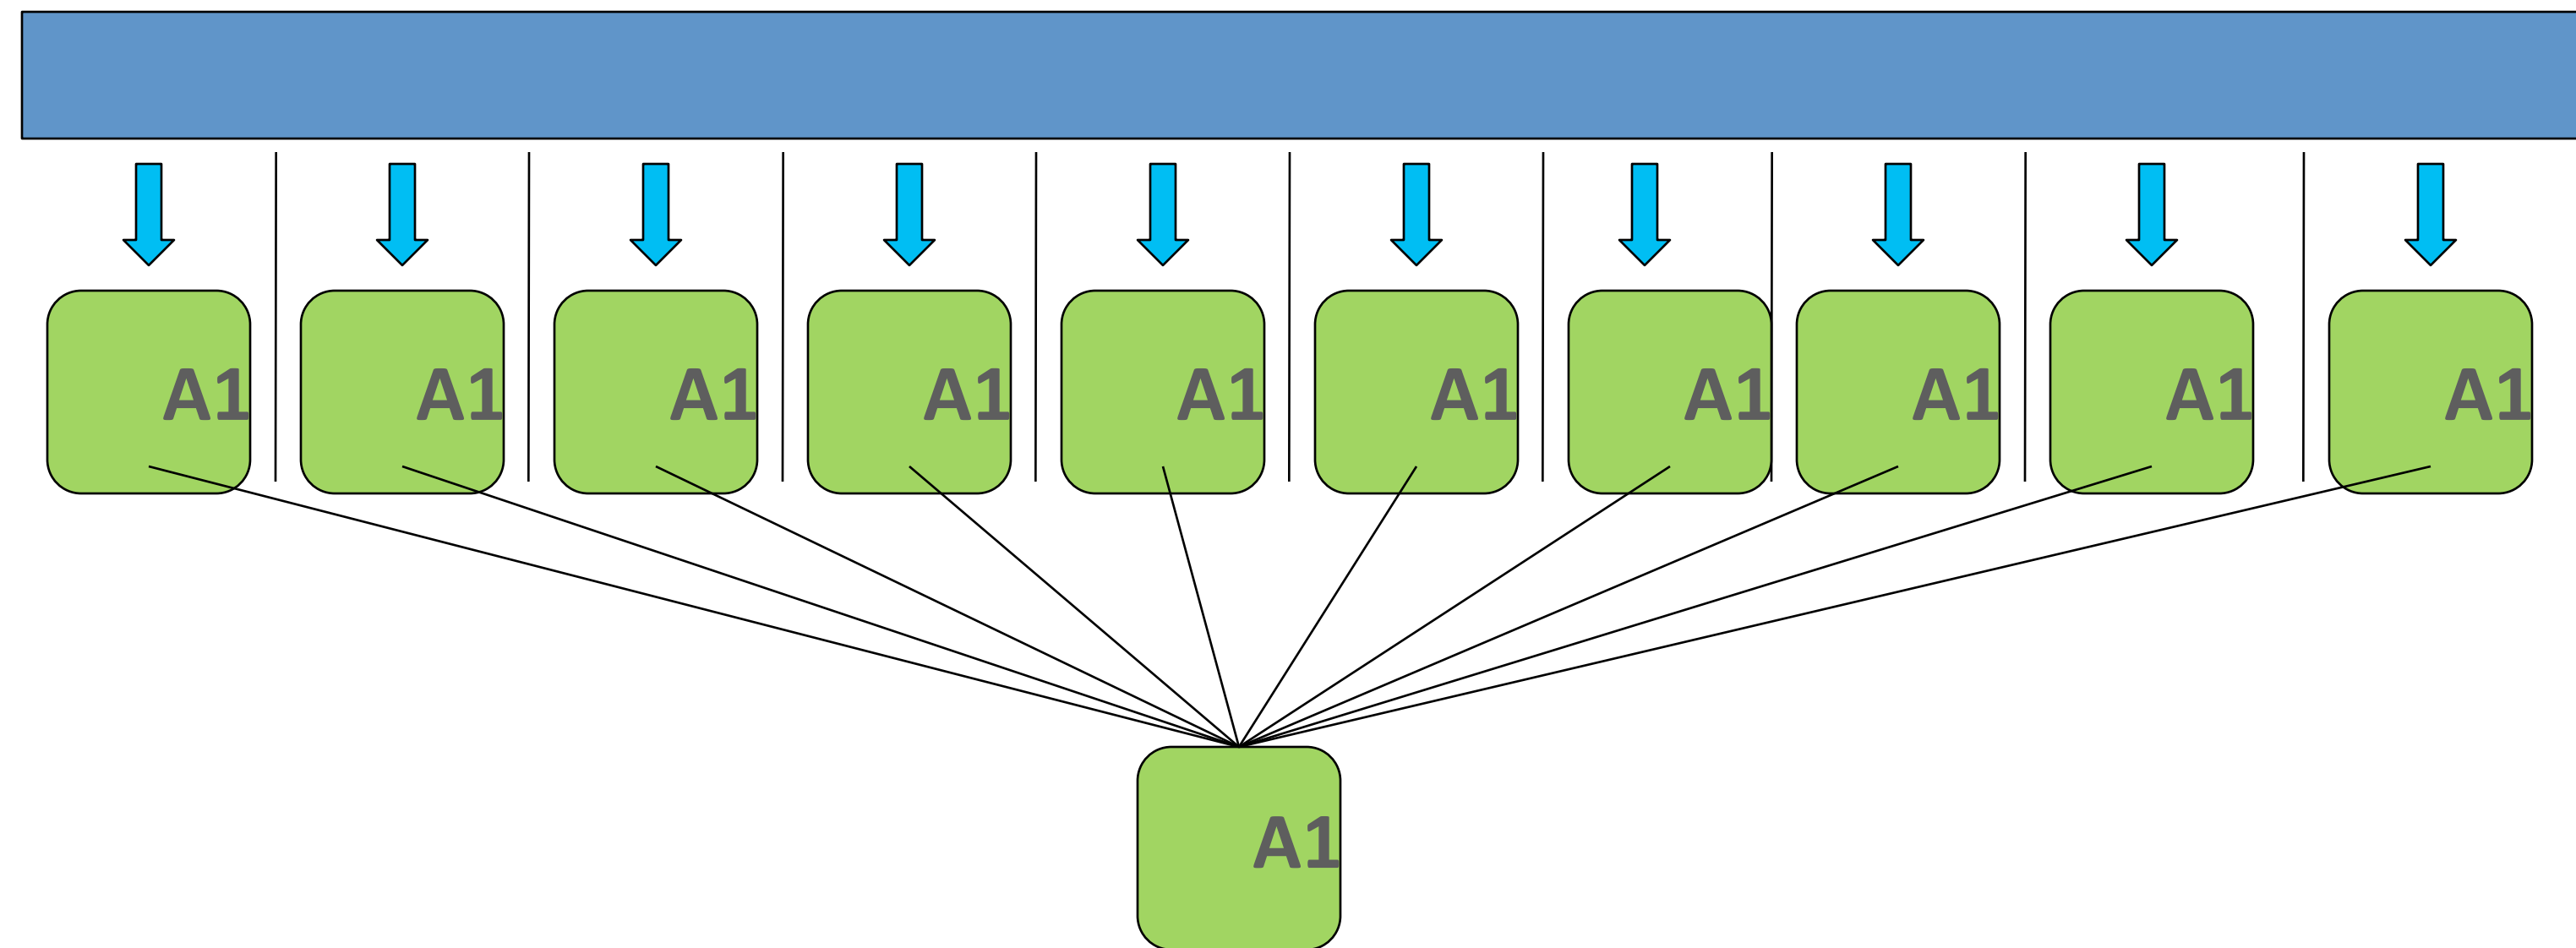
$\text{min} = A[i]$

Not optimal: $O(n^2)$ work

Optimal Min-find

- Balanced Binary tree
 - $O(\log n)$ time
 - $O(n)$ work => Optimal
- Use Accelerated cascading
- Make the tree branch much faster
 - Number of children of node $u = \sqrt{n_u}$
 - Where n_u is the number of leaves in u 's subtree
 - Works if the operation at each node can be performed in $O(1)$

From n^2 processors to $n\sqrt{n}$



Step 1: Partition into disjoint blocks of size \sqrt{n}

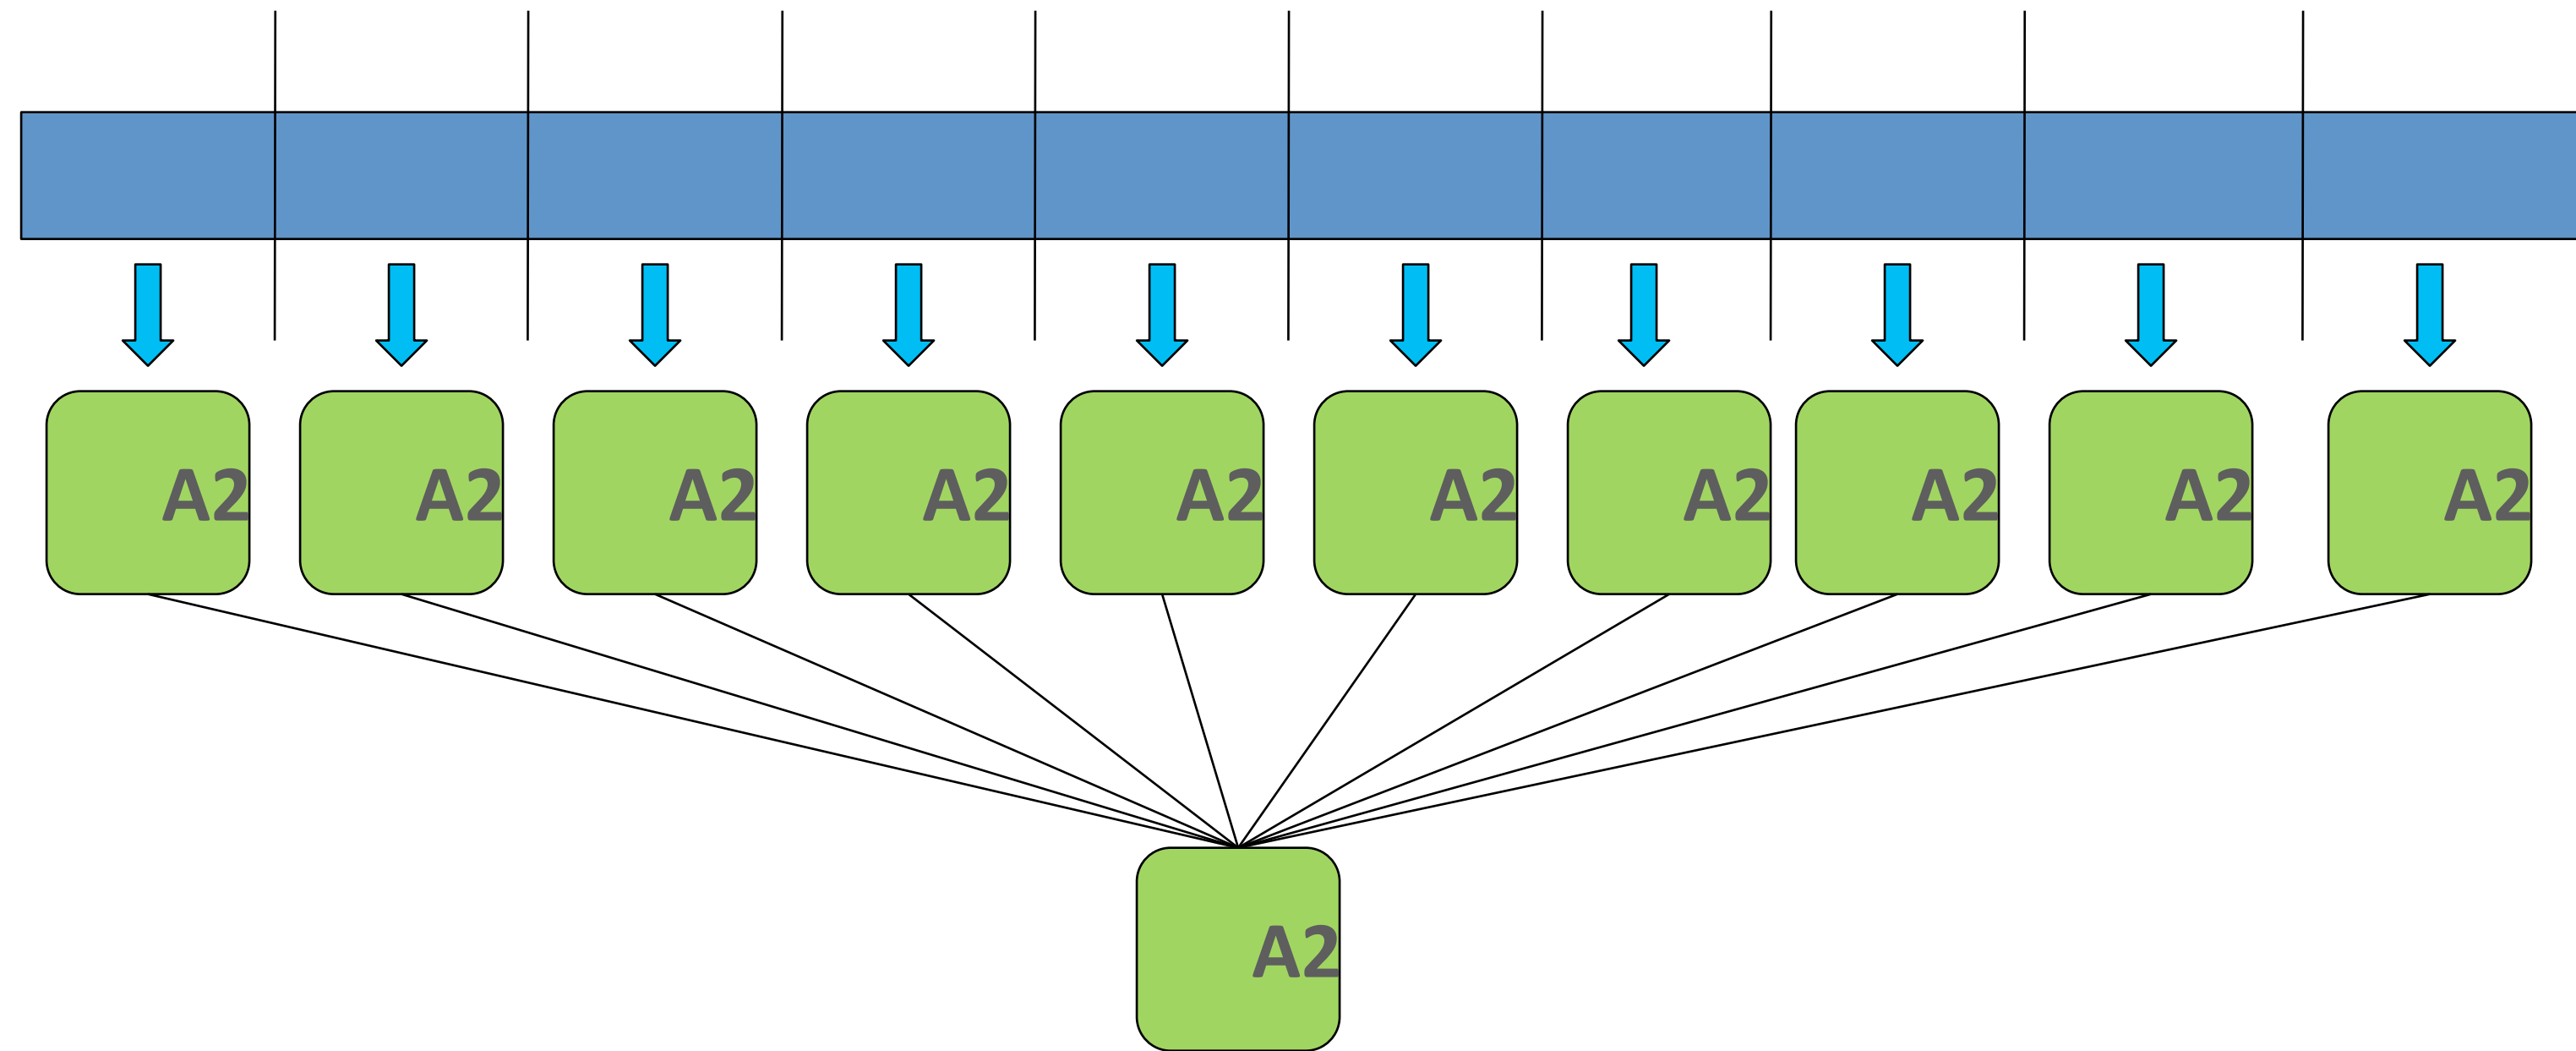
Step 2: Apply $A1$ to each block

Step 3: Apply $A1$ to the results from the step 2

$$n\sqrt{n}$$

$$n$$

From $n\sqrt{n}$ processors to $n^{1+1/4}$



Step 1: Partition into disjoint blocks of size \sqrt{n}

Step 2: Apply A_2 to each block


Step 3: Apply A_2 to the results from the step 2

$$n^{1/2} \cdot n^{3/4}$$

$$n^{3/4}$$

Algorithm A_{k+1}

1. Partition input array C (size n) into disjoint blocks of size $n^{1/2}$ each
2. Solve for each block in parallel using algorithm A_k
3. Re-apply A_k to the results of step 2: $n/n^{1/2}$ minima



$$\begin{array}{ccccccc}
 A_1 & & A_2 & & A_3 & & \dots \\
 n^2 & \rightarrow & n^{1+1/2} & \rightarrow & n^{1+1/4} & \rightarrow & n^{1+1/8} \rightarrow n^{1+1/2^k} \dots \sim n^{1+\epsilon}
 \end{array}$$

Algorithm A_∞ takes ?? with $n^{1+\epsilon}$ processors

Doubly logarithmic-depth tree

$n \log \log n$ work, $\log \log n$ time


$$n^{\frac{1}{2^i}} = O(1) \text{ at leaf}$$

Min-Find Review

- Constant-time algorithm
 - $O(n^2)$ work
- $O(\log n)$ Balanced Tree Approach
 - $O(n)$ work (Work-Optimal)
- $O(\log \log n)$ Doubly-log depth tree Approach
 - $O(n \log \log n)$ work
 - Degree is high at the root, reduces going down
 - #Children of node $u = \sqrt{\text{#nodes in tree rooted at } u}$
 - Depth = $O(\log \log n)$

Accelerated Cascading

- Solve recursively
- Start bottom-up with the optimal algorithm
 - until the problem sizes is smaller
- Switch to fast (non-optimal algorithm)
 - A few small problems solved fast but non-work-optimally
- Min Find:
 - Optimal algorithm for lower $\log \log n$ levels
 - Then switch to $O(n \log \log n)$ -work algorithm



n work, $\log \log n$ time